

232 req/s on one GPU; per-request work independent of tenant count

