

---

# User as Engram: Internalizing Per-User Memory as Local Parametric Edits

---

Bojie Li  
Pine AI

## Abstract

Personal memory in LLMs is two problems, not one: *content* (per-user facts) and *meta-skill* (reasoning patterns that turn facts into answers). Per-user LoRA conflates them and pays an architectural price: on Mini-Engram-d20 ( $n=20$  users, 3 seeds), it raises validation bpb on text *unrelated* to the user’s facts by  $\Delta=+1.78$ , while per-user Engram-row insertion raises it by  $+0.00005$ — **$\sim 34,000\times$  less, by construction**. This architectural gap is base-independent; whether it produces user-visible damage is not. Per-user LoRA hurts 85% of users on Mini-Engram-d20 (a base LM) but helps on every instruction-tuned base we tested (Qwen2.5-3B, Llama-3.1-8B, Mistral-7B-v0.3, Qwen2.5-7B; mean  $\Delta=+0.088$  to  $+0.217$ , 0–20% worse).

This motivates a **layered architecture**: one shared LoRA holds cross-user meta-skill (amortised over the population), per-user Engram-row overrides hold per-user content (local, 88 KB/user). At  $n=20$  test users, the layered design matches per-user LoRA’s direct recall (100%) at  $7.4\times$  higher indirect-reasoning accuracy (44% vs 6%),  $4.6\times$  less contamination ( $\Delta\text{bpb} +0.39$  vs  $+1.78$ ), and *never* hurts indirect reasoning relative to the no-adaptor base (0/60 vs 49/60 users pooled across seeds). The naive composition (per-user LoRA + per-user Engram) inherits LoRA’s damage (50/60 worse).

Against RAG, the ranking flips with KB size. At 34 facts/user, Qwen2.5-3B-Instruct + RAG top-3 wins (52% vs F’s 44% at zero context). At realistic KB ( $N \in \{100, \dots, 1000\}$  facts/user via a 30-user distractor pool), all-MiniLM-L6-v2 top-3 recall collapses  $7\times$  (62%  $\rightarrow$  9%) and Qwen-3B + RAG drops to **30%—14 pp below F** on a  $\sim 2.5\times$  larger backbone. F is invariant because the per-user Engram table does not grow with population size; substrate ranking is decided by deployment scale.

We release four Mini-Engram weights (178 M–1.22 B parameters; the first publicly available Engrams outside DeepSeek-AI), all training code, and complete result JSONs. A 50-line multi-tenant server reaches **232 req/s** on one GPU with zero cross-user leak by construction.

## 1 Introduction

Modern transformer (?) LLMs serve millions of users from a single backbone, but per-user *personal memory*—the system’s knowledge of *your* doctor, *your* preferences, *your* calendar—remains an open architectural problem. Three families address it: in-context learning (ICL, ??) prefixes the facts; retrieval-augmented generation (RAG, ??) fetches them at query time; and per-user adapters internalise them as weight deltas. The third family has accumulated a long list of variants on the same recipe—PRAG (?), DyPRAG (?), DistilledPRAG (?), OPPU (?), HYDRA (?), MemLoRA (?), T2L (?), and PER-PCS (?) all train a per-user rank- $\sim 64$  LoRA (??) on  $Q/K/V/O$  + MLP via NTP on an (observation, fact, QA) mixture, then apply it on top of a frozen base. We call this canonical

recipe **POLAR-class** and use it as our reference per-user-LoRA baseline throughout (rank-64 on  $Q/K/V/O$  in every layer; 12 epochs  $\times$  200 NTP samples per user; frozen base).

**The mechanism behind contamination.** A POLAR-class LoRA places a low-rank delta on  $Q/K/V/O$  in *every* transformer layer; that delta enters every forward pass, including ones whose tokens have nothing to do with the user’s facts. Two consequences decouple at measurement time. The *architectural* signature—validation bpb on unrelated text—is base-independent and large. The *user-visible* reasoning effect is base-dependent: instruction-tuned bases absorb the perturbation; base LMs do not. The load-bearing claim is the architectural one (Section ??).

**The fix: local per-user edits.** An edit avoids contamination iff it fires only when the trigger N-gram is present and is mathematically inert everywhere else. We propose **User as Engram**: surgical writes to the hash-keyed embedding table of an Engram-pretrained model (?), whose gated lookups fire only on suffix N-grams via deterministic multiplicative-XOR hashing. Per user, we write one fact-row per trigger N-gram into an *override map* applied per request. Locality is a measurable invariant, not an asymptotic approximation: on Mini-Engram-d12, a UNEMBED\_P insertion at the last Engram layer perturbs the residual stream by *exactly* 0.000 at every non-trigger position, through every subsequent attention+MLP layer (Section ??, Figure ??).

## Contributions.

1. **A measured  $34,000\times$  contamination gap** (Section ??). On Mini-Engram-d20, per-user LoRA produces  $\Delta\text{bpb} +1.78$  on unrelated text vs.  $+0.00005$  for per-user Engram-row insertion. The architectural gap is base-independent; the user-visible reasoning effect is not (85% of users worse on a base LM, 0–20% on four instruction-tuned bases).
2. **User as Engram: a local-edit substrate** (Sections ??–??). Surgical row-insertion via three strategies (UNEMBED\_P closed-form; OPT independent; Joint OPT). Joint OPT roughly doubles top-1 recall (68% vs. 36%) over independent OPT at 100 facts/user at identical 88 KB storage. We release four Mini-Engram checkpoints (178 M–1.22 B; the first publicly available Engrams outside DeepSeek-AI) and a 50-line multi-tenant server reaching **232 req/s** with zero cross-user leak.
3. **Capacity, fact-count, and dense-size scaling** (Sections ??–??). A 30-cell  $5\times 3\times 2$  ablation pins the capacity optimum; LOCOMO Joint OPT multi-token F1 scales from 0.159 to 0.310 between 178 M and 1.22 B. By category, Engram is  $+49\%/+178\%/-53\%$  on reasoning/multi-hop/open-domain (multi-hop wins come from suffix-N-gram surface overlap, not chained reasoning).
4. **A layered architecture that Pareto-dominates every per-user single-substrate baseline** (Section ??). Per-user Engram (content) + one shared LoRA (meta-skill) reaches 100% direct top-1,  $7.4\times$  higher indirect reasoning than per-user LoRA,  $4.6\times$  less contamination, and 0/20 (vs 17/20) users hurt on indirect, replicated across 3 seeds. The lead survives a personal  $\rightarrow$  medical schema shift ( $7.6\times$  over B with identical locality), unlike the recite-then-reason recipe which collapsed cross-schema (Section ??). The naive composition (per-user LoRA + Engram) is worst-of-both; rank ablation pins  $r=16$ .
5. **A KB-scale head-to-head with RAG** (Sections ??–??). At 34 facts/user, Qwen-3B + RAG wins (52% vs F’s 44%). At 1000 facts/user, retrieval recall@3 collapses  $7\times$  ( $62\% \rightarrow 9\%$ ) and Qwen-3B + RAG drops to 30%, **14 pp below F**, which is invariant because the per-user Engram table does not grow with population size.

**Paper roadmap.** Section ?? establishes the negative result (per-user LoRA contaminates). Section ?? introduces User as Engram and verifies locality. Section ?? releases four Mini-Engram checkpoints. Section ?? reports recall, density, storage, scaling, and head-to-head substrate comparisons. Section ?? reports the multi-tenant serving system. Section ?? introduces the layered architecture and its head-to-head with RAG across KB sizes. Section ?? reads off substrate-selection rules.

**Reproducibility.** Code, all four trained Mini-Engram checkpoints, and the raw result JSON files for every table and figure are available at <https://github.com/bojieli/user-as-engram>. Each major experiment is one bash script and reproduces in 1–48 h on a single Blackwell GPU (Appendix ??).

Figure 1: LoRA-as-memory is reasoning-negative on indirect questions.




Figure 1: The original  $n=10$  pilot (Section ??, Qwen2.5-3B-Instruct). **(a)** Per-user POLAR LoRA hits perfect direct recall (1.000); indirect recall (0.150) is below the no-adaptor base (0.170). **(b)** 5 of 10 users worse on indirect. **This pilot does not replicate at scale:** at  $n=30$  the mean delta is  $+0.088$  with  $6/30$  ( $\sim 20\%$ ) worse, and on three further instruction-tuned bases (Llama-8B, Mistral-7B, Qwen-7B at  $n=20$ ) the worse-fraction is 0–5%. The *architectural* contamination is robust regardless (Table ??): val\_bpb  $\Delta$  on Mini-Engram-d20 is  $+1.70$  (3-seed mean) for per-user LoRA vs  $+0.00005$  for Engram-row insertion.




Figure 2: Cost-quality Pareto across personal-memory substrates at the same answering LM (Mini-Engram-d20@1536). User-as-Engram Joint OPT matches a POLAR-class LoRA’s LOCOMO F1 at  $161\times$  smaller per-user storage (88 KB vs. 14.2 MB at 100 facts). Retrieval baselines sit on the steep low-storage side of the curve but cap out below Engram on quality.

## 2 Background

**The Engram architecture.** ? introduce *conditional memory* as a sparsity axis complementary to mixture-of-experts (MoE, ??). At each token position  $t$ , suffix N-grams of canonicalised tokens are hashed via  $K$  multiplicative-XOR heads into prime-sized embedding tables  $E_{n,k}$ . Retrieved rows  $e_t$  are projected by learned  $W_K, W_V$  and gated by an attention-style scalar  $\alpha_t = \sigma(\text{RMS}(h_t)^\top \text{RMS}(W_K e_t) / \sqrt{d})$ . The output  $\alpha_t W_V e_t$  is added to the residual stream at the selected Engram layer. Crucially, the address  $z_{t,n,k}$  is deterministic from the input token IDs—known before the forward pass—so the table can be offloaded to host DRAM without GPU contention.

**LoRA-as-memory family.** Per-user/per-document LoRA adapters (?????) train a low-rank delta (???) on  $Q/K/V$  (and sometimes MLP) matrices. The base is frozen; the LoRA is trained per user via NTP on a (observation, fact, QA) mixture (POLAR-class), via knowledge-distillation from a

Table 1: Per-user edit contamination on held-out text, Mini-Engram-d20 (canonical seed S0;  $n=20$  users).  $\Delta\text{bpb}$  of 0 means the edit leaves the base’s behavior on unrelated text mathematically unchanged. The 3-seed mean (S0, S1, S2) is reported in the last column and per-seed values are in Appendix ??; the contamination ratio ( $\sim 30,000\text{--}35,000\times$ ) is seed-stable.

| edit method           | S0 $\Delta\text{bpb}$ | S0 per-user range    | worse vs base (S0) | 3-seed mean $\Delta\text{bpb}$ |
|-----------------------|-----------------------|----------------------|--------------------|--------------------------------|
| no edit (baseline)    | 0.0000                | —                    | 0/20               | 0.0000                         |
| per-user LoRA $r=64$  | <b>+1.784</b>         | [+0.44, +3.74]       | <b>17/20 = 85%</b> | +1.698 [+1.56, +1.78]          |
| per-user Engram J-OPT | <b>+0.00005</b>       | [+0.00004, +0.00007] | 0/20               | +0.00005                       |

teacher (DyPRAG (?), DistilledPRAG (?)), or via hypernetwork emission (T2L (?)). All variants *edit the model weights globally*: every forward pass sees the LoRA delta.

**Memory architectures and adjacent work.** Engram (?) sits in a lineage of trainable key–value memory modules: PKM (?), PEER (?), Memory Layers at Scale (?), Ultra-Mem (?), OverEncoding (?), SCONE (?), BLT (?), and SuperBPE (?). None addresses surgical per-user insertion at inference. The closest analog in the diffusion-model world is LoRA stacking for Stable Diffusion (?); analogous mechanisms in the LLM world include LoRAHub (?) which trains a combiner network. Engram override maps with disjoint addresses are *additive without a combiner*, as we show in Section ??.

**Personal memory benchmarks.** Long-term personal memory is evaluated by LOCOMO (?), LongMemEval (?), BEAM (?) (multi-turn evolving beliefs), and PersonaMem-v2 (?). We use synthetic per-user fact corpora for the within-user fact-density sweeps (tighter control over the parameter we vary) and LOCOMO (Section ??) for end-to-end evaluation against retrieval baselines. We chose LOCOMO over LongMemEval and BEAM because its per-question evidence pointers let us train a per-fact Engram row and a retrieval entry against the same gold sentence—a fair apples-to-apples “same gold, different substrate” comparison; LongMemEval and BEAM remain open for future work.

**Knowledge editing and recitation.** Knowledge editing methods (ROME (?), MEMIT (?)) edit FFN weights via causal-traced locations; ripple-effect benchmarks like MQuAKE (?), RippleEdits (?), and CounterFact (?) measure indirect consequences. Recitation methods (SR-RAG (?), Hewitt et al. (?), “Thinking to Recall” (?), recitation-augmented generation (?)) train or prompt the model to surface stored facts before reasoning. We apply the recitation idea to per-user adapters in Section ?? (within-schema 0.41, but cross-schema collapses to 0.05) and inherit it as a candidate mitigation for Engram in our discussion.

### 3 Per-user LoRA contaminates globally

A per-user LoRA edits  $Q/K/V/O$  projections in every transformer layer; the edit is global by construction. We measure its two properties directly: (i) the *architectural* contamination (validation bpb on text unrelated to the user’s facts), which is predictable and large; (ii) the *user-visible* reasoning effect, which is variable and depends on whether the base is instruction-tuned.

#### 3.1 Architectural contamination on the same base (Mini-Engram-d20)

We hold the base fixed (Mini-Engram-d20@1536, 1.22 B parameters) and train a per-user POLAR-class LoRA (rank-64 on  $Q/K/V$ , 1,500 steps per fact, single-token gold) on the synthetic user fact sets used throughout. For each of 20 test users we measure  $\text{val\_bpb}$  on a held-out 262 K-token ClimbMix shard—text unrelated to the user’s facts—before and after the edit. Per-user Engram-row Joint OPT is measured under the identical protocol on the same base.

The ratio is  $\sim 34,000\times$  on S0 and  $\sim 33,000\times$  on the 3-seed mean. Per-user LoRA *more than triples*  $\text{val\_bpb}$  on unrelated text (mean  $0.74 \rightarrow 2.52$  on S0); per-user Engram-row insertion leaves it unchanged to four decimal places in every seed. The mechanism is visible by construction: **LoRA’s edit is global; Engram’s edit fires only on the trigger N-gram’s hash addresses**. Pooled across 60

Table 2: Cross-base LoRA scaling. “ $\Delta$  user” is mean indirect recall with the adapter minus the same user’s recall without the adapter. “worse” is the fraction of users with strictly negative  $\Delta$ .

| base                                 | $n$       | mean $\Delta$ user-indirect | $\Delta$ range   | worse than base    |
|--------------------------------------|-----------|-----------------------------|------------------|--------------------|
| Qwen2.5-3B-Instruct (original pilot) | 10        | −0.008                      | —                | 6/10 = 60%         |
| Qwen2.5-3B-Instruct (scaled)         | <b>30</b> | <b>+0.088</b>               | [−0.143, +0.214] | <b>6/30 = 20%</b>  |
| Llama-3.1-8B-Instruct                | 20        | <b>+0.188</b>               | [+0.030, +0.333] | <b>0/20 = 0%</b>   |
| Mistral-7B-Instruct-v0.3             | 20        | <b>+0.217</b>               | [+0.091, +0.344] | <b>0/20 = 0%</b>   |
| Qwen2.5-7B-Instruct                  | 20        | <b>+0.132</b>               | [−0.030, +0.242] | <b>1/20 = 5%</b>   |
| Mini-Engram-d20 (base LM)            | 20        | −0.133                      | [−0.300, +0.050] | <b>17/20 = 85%</b> |

user-experiments (3 seeds), per-user Engram J-OPT makes 0/60 users worse than no-edit on indirect reasoning; per-user LoRA makes 49/60 (82%) worse.

### 3.2 User-visible effect varies with the base model

We replicate the POLAR per-user LoRA recipe (rank-64 NTP on observation/fact/QA mixtures, 12 epochs  $\times$  200 samples) on three instruction-tuned bases and one base LM (Table ??):

The pattern splits cleanly. On the **base LM** (Mini-Engram-d20), LoRA’s global perturbation disrupts fragile completion behaviour: 17/20 worse, mean  $\Delta = -0.133$ . On all four **instruction-tuned bases**, the base’s reasoning skill absorbs the perturbation: mean  $\Delta$  from +0.088 to +0.217, worse-fraction 0–20%. The original POLAR pilot’s “5/10 worse than base” headline was a high-variance  $n=10$  artifact: scaling Qwen-3B to  $n=30$  drops the worse-fraction to 20%; Llama-3.1-8B and Mistral-7B reach 0%.

### 3.3 Architectural lesson

**Any unconditionally active edit contaminates unrelated forwards**; conditioning on the trigger keeps the base mathematically unchanged elsewhere. Two recipe-level fixes do not recover this property. *Recite-then-reason traces* mixed into Stage A close the within-schema indirect gap on Qwen-3B (0.41) but collapse cross-schema (0.046, worse than base). A full-parameter meta-train over the LoRA distribution (*Stage C*) lifts training-user indirect by +0.122 but transfers +0.000 to held-out users, and the base fine-tune misaligns frozen LoRAs enough to collapse direct recall to 0.22 (Appendix ??). The fix must be architectural: Section ?? introduces it; Section ?? restores the meta-skill that per-user LoRA was attempting to provide.

## 4 Method: User as Engram

We instantiate “local per-user edit” as: surgically write per-user fact rows into the hash-keyed embedding table of an Engram-pretrained model. Figure ?? shows the serving architecture; the rest of this section walks through each component.

### 4.1 Surgical row insertion

Tokenise the trigger to  $x_1, \dots, x_T$ . The trigger position is  $t^* = T - 1$ . For each  $(n, k)$  in the configured Engram layer, the addresses  $\{z_{t^*, n, k}\}$  are deterministic from the trigger tokens. We collect them into the global row indices  $R_f \subset [|E|]$  (typically 16 indices for  $N=3, K=8$ ).

**Three insertion strategies** (in increasing cost and quality; Appendix Figure ?? plots their recall side by side against a RANDOM control):

- **UNEMBED\_P** (closed-form): solve  $W_V e^* \approx U_y$  via the Moore–Penrose pseudo-inverse  $e^* = W_V^\dagger U_y$ , where  $U_y$  is the unembed direction of the answer’s first token. Write  $e^*$  into all rows in  $R_f$ . Cost: 1 mat-vec, < 1 ms per fact.
- **OPT independent**: initialise from UNEMBED\_P, then 15 Adam steps on  $e$  to maximise  $\log p(y \mid \text{trigger})$ . Cost: 15 fwd+bwd,  $\sim 1$  s per fact. Each fact optimised in isolation.




Figure 3: EngramServer architecture. Per user, an override map of (row\_idx, row\_vector) pairs lives in DRAM. On each request, the server saves the originals at the affected addresses ( $\sim 2$  ms), writes the user’s overrides, runs the forward pass, then restores. The Engram lookup at the configured layer transparently sees the override values; the gate fires only at the trigger N-gram. There is no router and no graph rewrite.




Figure 4: Engram surgical insertion is spatially perfect. Heatmap of  $\|x_{\text{after}}^{(\ell)} - x_{\text{before}}^{(\ell)}\|$  for a UNEMBED\_P insertion at the last Engram layer ( $L_{\text{eng}}=7$ , Mini-Engram-d12). The diff is exactly 0.000 at every position before  $L_{\text{eng}}$  (causality) and at every non-trigger position after  $L_{\text{eng}}$ , through 5 subsequent attention/MLP layers. Compare with LoRA, where every position in every forward sees the global delta.

- **Joint OPT** (the recommended default for  $\geq 30$  facts/user): allocate one trainable row tensor over the union of all touched addresses across the user’s facts. Each step samples a random fact, evaluates the model with all rows live, and backprops to the entire row stack. By coordinating rows during training, Joint OPT eliminates the cross-fact forward-pass interference that limits independent OPT under high density (Section ??).

## 4.2 Local edit verified

Writing a UNEMBED\_P marker at the last Engram layer and measuring  $\|x_{\text{after}}^{(\ell)} - x_{\text{before}}^{(\ell)}\|$  at every layer and every position yields the property by inspection (Figure ??): the diff is *exactly* 0.000 at every non-trigger position through 5 subsequent attention/MLP layers, while the trigger-position diff is 123.0 at layer 8 and 1.54 at the final layer. **This is the architectural property that distinguishes Engram from LoRA:** no contamination of non-trigger forwards.

Table 3: Mini-Engram pretraining at the ablation-optimal recipe. All use the large Engram ( $50\text{ K} \times 256 = 51.2\text{ M}$  Engram params). Trained on a single NVIDIA RTX PRO 6000 Blackwell (102 GB) in bf16.

|                          | d8 v2          | d12 v2          | <b>d12@1280 opt.</b> | <b>d20@1536 opt.</b>     |
|--------------------------|----------------|-----------------|----------------------|--------------------------|
| depth $\times$ width     | $8 \times 512$ | $12 \times 768$ | $12 \times 1280$     | $20 \times 1536$         |
| Engram layers            | 2, 5           | 2, 7            | 2, 7                 | 2, 11                    |
| total params             | 178 M          | 339 M           | 625 M                | <b>1.22 B</b>            |
| scaling-params           | 42 M           | 110 M           | 278 M                | 617 M                    |
| ClimbMix tokens (12 t/p) | 0.50 B         | 1.32 B          | 3.34 B               | 7.40 B                   |
| final val_bpb            | 0.912          | 0.827           | 0.770                | <b>0.730</b>             |
| wall-clock (single GPU)  | $\sim 50$ min  | $\sim 5$ h      | $\sim 10.5$ h        | $\sim 70$ h (shared GPU) |

### 4.3 Per-user override tables and additive composition

The production design is per-user override tables. Each user has a small dictionary  $\{r_i \mapsto v_i\}$  of fact rows. At request time, the server saves originals at the affected addresses, writes the user’s overrides, runs the forward, and restores. Cross-user leakage is *zero by construction*—user A’s overrides are not in the table when user B queries.

Override maps with disjoint addresses commute, so corporate facts + user facts (or any number of domain-specific Engrams) **stack additively at inference** without retraining a combiner. This mirrors how Stable Diffusion LoRAs additively stack, but at the row level. Section ?? confirms additive composition is lossless when the domains have disjoint trigger templates.

## 5 Nano-scale public reproduction of Engram

? released only architectural reference code; no Engram-trained weights are public. We graft Engram into Karpathy’s nanochat (?) (a GPT-2-style backbone (?) optimised with Muon + AdamW (??)) and train four Mini-Engram models on a single Blackwell GPU at the ablation-optimal recipe (Section ??): the same large Engram table ( $50\text{ K} \times 256$ , 51.2 M params) at Karpathy 12 t/p of scaling-params for every dense size.

**Model-name notation.** Throughout the paper we use  $d\{\text{depth}\}@ \{\text{width}\}$  for the ablation-optimal Mini-Engrams (d12@768, d12@1280, d20@1536) and v1/v2 to denote earlier and final training runs at the same shape. Where the width is omitted (e.g. d12 without a suffix), we mean the v2 (ablation-optimal at width 768, 339 M total) checkpoint. The four headline trained checkpoints and their token budgets are: d8 v2 (178 M total, 42 M scaling, 0.50 B tokens), d12 v2 a.k.a. d12@768 (339 M / 110 M / 1.32 B), d12@1280 (625 M / 278 M / 3.34 B), and d20@1536 (1.22 B / 617 M / 7.40 B). All four use the same 51.2 M Engram-table.

**Sensitivity asymmetry reproduced (?, their §6.3).** Suppressing Engram (engram\_suppress=True) costs +0.035 bpb on validation. Factual top-5 retains 93.3% of active performance; reading top-5 retains 100%. The asymmetry  $\Delta = +0.067$  matches the direction of ? Figure 6 (TriviaQA: 29% retained, C3: 93%,  $\Delta \approx +0.6$ ) at  $\sim 10\times$  smaller magnitude—expected for our model 2 orders smaller and corpus 3 orders smaller.

**Effective deepening reproduced (?, their §6.1).** LogitLens layer-3 KL is 3.66 *lower* for engram-d8 than base-d8, matching the shape of ? Figure 4(a) (Appendix Figure ??).

## 6 Experiments

This section reports five blocks of measurements, each answering a distinct question about User-as-Engram: **(a) per-fact recall and density** (Sections ??–??): how accurately do single and multiple per-user facts surface? **(b) cost** (Sections ??–??): storage and wall-time vs. POLAR-class LoRA. **(c) head-to-head with memory systems / RAG / LOCOMO** (Sections ??–??): how does the substrate

Table 4: Single-fact-at-a-time recall (each fact tested in isolation with restore between facts; XL corpus, 1000 facts).

| method                               | USER top-1   | ORG top-1    |
|--------------------------------------|--------------|--------------|
| baseline (no insertion)              | 1.7%         | 8.0%         |
| ICL (= optimal RAG)                  | 95.8%        | 99.9%        |
| UNEMBED_P                            | 11.2%        | 14.0%        |
| <b>User-as-Engram OPT</b>            | <b>87.6%</b> | <b>90.8%</b> |
| SFT-LoRA per-fact (rank 8, 30 steps) | 100%         | 100%         |

Figure 3: Joint OPT vs LoRA-rank-64 across fact-density per user.




Figure 5: Within-user fact-density: top-1 (left) and top-5 (right) recall vs. number of facts inserted simultaneously into one user’s override map. Joint OPT (blue) closes most of the gap to LoRA rank-64 (red dashed) at  $161\times$  less storage (88 KB vs. 14.2 MB at 100 facts).

compare against retrieval baselines on the same answering LM? **(d) scaling ablations** (Sections ??–??): capacity vs. token budget vs. dense-size. **(e) probes beyond direct recall** (Sections ??–??): multi-hop, free-continuation, paraphrase. We use three corpora throughout: the **base** 200-fact benchmark (100 USER + 100 ORG facts), the **XL** corpus of 1,000 USER + 1,000 ORG templated facts (Section ??), and the **XXL** corpus of 3,132 distinct trigger templates (high-density / distinct-template stress tests).

## 6.1 Single-fact recall

OPT achieves 87.6–90.8% top-1 recall on 1 000 single-fact tests, within  $\sim 8$  points of the ICL ceiling (95.8%) at three orders of magnitude lower per-user cost than POLAR-style LoRA training.

## 6.2 Within-user fact density and Joint OPT

Loading many facts simultaneously creates context-position interference among the live rows in one user’s override table. We sweep the density curve on Mini-Engram-d12 with the XXL corpus (Figure ??).

**Joint OPT is the recommended default for  $\geq 30$  facts/user.** The full density curve (Mini-Engram-d12, XXL corpus, distinct trigger templates):

**Top-5 stays  $\geq 90\%$  all the way to 300 facts/user.** At 100 facts, Joint OPT achieves 68% top-1 / 96% top-5 in 88 KB vs. multi-fact LoRA rank-64’s 100%/100% at 14.2 MB ( **$161\times$  more storage**). At 1 000 facts: Joint OPT 35%/72% in 296 KB vs. LoRA rank-64’s 38%/77% at 14.2 MB ( **$48\times$  more storage**).

**The density ceiling is set by the Engram table, not the dense backbone.** Repeating the sweep on the two larger Mini-Engrams at fixed Engram capacity (50 K  $\times$  256, 51.2 M Engram params; Table ??), top-5 improves slightly at  $n = 100$  (0.96  $\rightarrow$  0.99), but at  $n = 1000$  top-1 *degrades* from



Table 5: Joint OPT density curve at single-user scale (Mini-Engram-d12 v2).

| N facts        | 30    | 100   | 300   | 1000  |
|----------------|-------|-------|-------|-------|
| top-1          | 0.733 | 0.680 | 0.507 | 0.351 |
| top-5          | 0.900 | 0.960 | 0.923 | 0.715 |
| train wall (s) | 11    | 44    | 29    | 228   |
| storage (KB)   | 36    | 88    | 156   | 296   |

Table 6: Joint-OPT density (top-1 / top-5) across three dense scales at fixed Engram capacity (51.2 M params). Top-1 at  $n=1000$  degrades with dense scale (gradient interference dominates).

| $n$ facts   | d12@768 (339 M)    | d12@1280 (625 M) | d20@1536 (1.22 B) |
|-------------|--------------------|------------------|-------------------|
| 100         | 0.68 / 0.96        | 0.65 / 0.98      | 0.67 / 0.99       |
| 300         | 0.51 / 0.92        | 0.46 / 0.87      | 0.48 / 0.90       |
| <b>1000</b> | <b>0.35</b> / 0.71 | 0.28 / 0.67      | 0.28 / 0.66       |

0.35 (d12@768) to 0.28 (larger). When the table saturates with 1000 facts, gradient interference is the binding constraint, and bigger dense backbones compete more strongly with the inserted rows for the gold logit. This makes the recipe change in Section ?? (*multi-fact insertion in the loss*) the most promising path to break the ceiling.

### 6.3 Storage scaling at deployment

At realistic deployment scales ( $10^4$ – $10^6$  users) the  $\sim 161\times$  ratio is the difference between fitting on one server vs. requiring distributed storage.

### 6.4 Method comparison: ICL, RAG, SFT-LoRA, POLAR LoRA

Figure ?? plots these rows on a cost-quality Pareto frontier where the storage axis spans four orders of magnitude. The trade-off has two regimes. **Single-fact:** OPT-15 matches the ICL ceiling (100%) at 88 KB and three orders of magnitude lower per-user cost than POLAR-style training. **Multi-fact:** every LoRA rank  $\geq 8$  saturates recall, so the comparison collapses to storage. At 100 facts, Joint OPT trails LoRA by  $\sim 35$  pt top-1 / 2 pt top-5 at  $20\times$  less storage (88 KB vs. 1.8 MB rank-8). At 1,000 facts the gap inverts: Joint OPT 35% vs. LoRA-rank-64’s 38% top-1, at  $48\times$  less storage. LoRA’s storage premium buys headroom that has already been spent at the rank-8 ceiling, not asymptotic recall.

### 6.5 Multi-domain additive composition

Two override maps with disjoint addresses commute, so corporate and user Engrams stack additively at inference without retraining a combiner. We verify this with a 10-fact corporate override (Globex facts) plus a 10-fact user override (“my doctor” facts): address overlap is 2.7% and both domains retain their alone-recall under composition (corp 80%/90%  $\rightarrow$  80%/90%; user 90%/100%  $\rightarrow$  90%/100%). When domains *share* templates, pair-overlap rises and the later-applied override wins shared addresses (a 100-fact corp + 100-fact user stack: user 64% unchanged, corp drops 58%  $\rightarrow$  25%). Architectural rule: **additive composition for disjoint-template domains; per-user override tables for same-template tenants.**

### 6.6 Comparison against memory systems

We compare User-as-Engram against state-of-the-art memory systems that target the same use case (per-user / personal memory). All systems use the *same* Mini-Engram-d12 as the final-answer LM; the only difference is how facts are stored and retrieved. The sentence-encoder used by RAG / MEM0 / MEMMACHINE baselines is all-MiniLM-L6-v2 (80M parameters; comparable in size to a production retriever).

Figure 4: Storage scaling. Engram is 200-1700 $\times$  smaller than LoRA.




Figure 6: Per-user storage scales linearly with user count. Engram override is  $\sim 161\times$  smaller at 100 facts/user (88 KB vs. 14.2 MB) and the gap widens to  $\sim 1700\times$  at 10 facts/user (LoRA’s parameter count is fact-count-independent while Engram grows linearly). At 1 M users  $\times$  100 facts/user, Engram storage is 100 GB vs. POLAR LoRA’s 40 TB.

Table 7: Method comparison at 100 facts/user on the ablation-optimal Mini-Engram-d12@1280 (625 M). Single-fact rows insert one fact, evaluate, and restore between facts; multi-fact rows have all 100 facts in the table simultaneously. Storage is the per-user delta; LoRA storage assumes fp32.

| method                                                       | top-1       | top-5       | wall (100 facts) | storage      |
|--------------------------------------------------------------|-------------|-------------|------------------|--------------|
| <i>Single-fact (one at a time):</i>                          |             |             |                  |              |
| baseline (no insertion)                                      | 0%          | 7%          | 0                | 0            |
| ICL (= optimal RAG)                                          | 100%        | 100%        | —                | 100 KB ctx   |
| UNEMBED_P (closed-form)                                      | 5%          | 35%         | <0.1 s           | 88 KB        |
| OPT-15 independent                                           | <b>100%</b> | <b>100%</b> | $\sim 50$ s      | 88 KB        |
| <i>Multi-fact (100 simultaneous, the production regime):</i> |             |             |                  |              |
| <b>Joint OPT (ours)</b>                                      | <b>65%</b>  | <b>98%</b>  | <b>44 s</b>      | <b>88 KB</b> |
| Multi-fact LoRA rank 8 (2000 steps)                          | 100%        | 100%        | 81 s             | 1.8 MB       |
| Multi-fact LoRA rank 32                                      | 100%        | 100%        | 81 s             | 7.1 MB       |
| <b>Multi-fact LoRA rank 64 (POLAR-class)</b>                 | 100%        | 100%        | 81 s             | 14.2 MB      |
| Multi-fact LoRA rank 128                                     | 99%         | 100%        | 81 s             | 28.3 MB      |
| SFT-LoRA per-fact (rank 8, 30 steps/fact)                    | 100%        | 100%        | 65 s             | 1.7 MB       |

This subsection probes *direct* fact recall under trigger-exact and paraphrased queries; the orthogonal *indirect-reasoning* comparison (does retrieval recover the gold fact across hops, and what happens as the KB grows?) is in Sections ??–??.

**Setup.** 100 USER facts (XXL corpus). Two query distributions:

- **Trigger-exact** (Table ??): the query is the fact’s stored prompt (e.g. “My favorite spice is”). Easy retrieval setting—the query prefix matches the fact verbatim.
- **Paraphrased** (Table ??): the query is a different surface form (e.g. “I love the spice”, “My preferred spice is”) of the same underlying fact. 16 facts  $\times$  4 paraphrase queries each.

**In trigger-exact, RAG dominates trivially**, because the query’s surface form is identical to the fact’s prefix; nearest-neighbour retrieval on a sentence-encoder space is near-perfect. Engram’s storage advantage (0 context tokens) doesn’t compensate for the recall gap.

**The story flips on paraphrased queries.** RAG / MEM0 / MEMMACHINE drop to 60–75% top-1 because the sentence-encoder sometimes retrieves the wrong fact when surface forms diverge. Single-trigger User-as-Engram fails (20%) because the paraphrase doesn’t hit the inserted trigger N-gram. But multi-trigger insertion (insert at every anticipated paraphrase—the 4 paraphrases plus

Table 8: Trigger-exact queries (100 facts on the XXL corpus, query = fact prompt prefix). Each system uses the same Mini-Engram-d12 to generate the final answer. The XXL corpus differs from the Section ?? XL corpus in trigger-template count (3,132 distinct templates vs. 1,000 templated facts); baseline numbers therefore do not match Table ?? exactly.

| method                             | top-1        | top-5       | avg context tokens |
|------------------------------------|--------------|-------------|--------------------|
| NO_MEMORY                          | 1.0%         | 7.0%        | 0                  |
| MARKDOWN_ALL (full fact dump)      | 21.0%        | 48.0%       | 1124               |
| <b>RAG_TOP1</b> (sentence encoder) | <b>99.0%</b> | <b>100%</b> | 16                 |
| RAG_TOP3                           | 96.0%        | 100%        | 40                 |
| MEM0_LIKE (top-5 retrieval)        | 94.0%        | 100%        | 63                 |
| MEMMACHINE_LIKE (top-3 episodic)   | 93.0%        | 100%        | 46                 |
| User-as-Engram OPT independent     | 36.0%        | 54.0%       | <b>0</b>           |
| User-as-Engram Joint OPT           | 68.0%        | 96.0%       | <b>0</b>           |

Table 9: Paraphrased queries (16 facts  $\times$  4 paraphrases each; this is a distinct protocol from the 4-fact  $\times$  5-paraphrase Mini-Engram-d8 rank study in Appendix ??, which probes per-paraphrase rank rather than aggregate accuracy). Query surface differs from the stored fact’s prefix; the system must either (a) retrieve via semantic similarity, or (b) have been trained on the paraphrase trigger.

| method                                  | top-1        | top-5        |
|-----------------------------------------|--------------|--------------|
| NO_MEMORY                               | 6.3%         | 14.1%        |
| MARKDOWN_ALL (158 tokens, all 16 facts) | 60.9%        | 73.4%        |
| RAG_TOP1                                | 64.1%        | 87.5%        |
| RAG_TOP3                                | 65.6%        | 82.8%        |
| MEM0_LIKE (top-5)                       | 62.5%        | 81.3%        |
| <b>MEMMACHINE_LIKE</b> (top-3)          | <b>75.0%</b> | <b>87.5%</b> |
| User-as-Engram OPT single-trigger       | 20.3%        | 25.0%        |
| <b>User-as-Engram OPT multi-trigger</b> | <b>96.9%</b> | <b>98.4%</b> |

the canonical trigger, 5 triggers per fact,  $\sim$ 5 s per fact) achieves **96.9% top-1**, vs. MEMMACHINE’s **75.0%**—a 22-point lead at zero context cost.

**Scale invariance.** We repeated both the trigger-exact and paraphrase comparisons on the optimal d12@1280 (625 M) and d20@1536 (1.22 B) Mini-Engrams. The shape is identical: retrieval baselines hit 1.00 top-1 on trigger-exact at d12@1280 and d20@1536; User-as-Engram multi-trigger holds 0.94–0.95 paraphrase top-1 vs. the best retrieval baseline at 0.81–0.83 (best is RAG\_TOP3 at d12@1280, MEMMACHINE at d20). The verdict—*retrieval wins when the query is verbatim the stored prefix, multi-trigger Engram wins under paraphrase*—does not change with dense scale. Engram numbers are essentially fixed; only the retrieval-baseline numbers move slightly because their answering-LM quality improves.

**Limitations of this comparison.** Our MEMMACHINE\_LIKE baseline is the retrieval primitive (top-3 sentence-encoder retrieval over atomic facts), not the full MemMachine system on conversational episodes (LOCOMO (?), LongMemEval (?)); the canonical MEM0 evaluation includes an LLM-based fact-extraction step at insert time, which we skip because our facts are already structured. We do not test queries for facts that were *never inserted* (the correct behaviour is “no information available”). The LOCOMO follow-up that addresses the multi-turn-dialogue regime is Section ??.

## 6.7 LOCOMO single-hop fact recall

To evaluate the same memory systems on a published long-term conversational benchmark, we run all eight systems from Section ?? on LOCOMO (?). We use *all 10 LOCOMO conversations* and take the first 80 single-hop (non-adversarial) questions per conversation, for a total of  $\sim$ 800 question-answer pairs grounded in dialog per model.

Table 10: LOCOMO single-hop, full 10 conversations ( $\sim 800$  QAs per model), *token-F1* metric. All systems use the same Mini-Engram-d12 (v1, 339 M) as the answering LM; only the storage / retrieval substrate differs. The corresponding LLM-judge results in Table ?? flip the ranking.

| method                                 | token-F1     |
|----------------------------------------|--------------|
| NO_MEMORY (no context)                 | 0.036        |
| MARKDOWN_ALL (full evidence in prompt) | 0.075        |
| RAG_TOP1 (sentence-encoder)            | 0.101        |
| RAG_TOP3                               | 0.117        |
| MEM0_LIKE (top-5 retrieval)            | 0.131        |
| MEMMACHINE_LIKE (top-3 episodic)       | 0.116        |
| User-as-Engram OPT (per-fact)          | 0.127        |
| <b>User-as-Engram Joint OPT</b>        | <b>0.169</b> |

Table 11: LOCOMO Joint OPT *token-F1* vs. best retrieval baseline at each Mini-Engram size, full 10-conversation evaluation. Engram appears to widen its lead with scale—but Table ?? shows token-F1 over-credits Engram.

| answering LM                     | params        | best retr. F1 | Engram J-OPT | token-F1 lead |
|----------------------------------|---------------|---------------|--------------|---------------|
| Mini-Engram-d8 v2                | 178 M         | 0.088         | 0.134        | +52%          |
| Mini-Engram-d12 v2               | 339 M         | 0.131         | 0.169        | +29%          |
| Mini-Engram-d12@1280 opt.        | 625 M         | 0.160         | 0.176        | +10%          |
| <b>Mini-Engram-d20@1536 opt.</b> | <b>1.22 B</b> | <b>0.169</b>  | <b>0.233</b> | <b>+38%</b>   |

**Setup.** Per LOCOMO question, the gold “evidence” field points to dialog turns containing the answer; we use the text of those turns as the unit stored by the retrieval baselines (RAG, MEM0, MEMMACHINE). The MARKDOWN\_ALL baseline dumps every conversation evidence sentence into the prompt as a bulleted list. For User-as-Engram, we extract a (question, first-token of answer) pair per QA and OPT-train it either independently (per-fact insertion) or jointly (sparse-row joint OPT, 2000 shared steps). All systems use the same Mini-Engram-d12 as the answering LM, which generates 16 tokens greedily; we score token-F1 against the gold answer with simple whitespace tokenization, lowercased, with punctuation stripped.

**User-as-Engram Joint OPT reaches token-F1 0.169 vs. 0.131 for the strongest retrieval baseline** (MEM0\_LIKE, +29% relative), and per-fact independent OPT also clears every retrieval baseline (0.127). Absolute scores are low because Mini-Engram-d12 v2 is a 339 M base LM with no instruction tuning and token-F1 punishes verbosity; the relative ordering is what matters. This is a best-case-each-system benchmark: retrieval stores the gold evidence sentence, Engram stores the gold  $(q, a)$  pair; the question isolated is which substrate extracts the answer most reliably under a fixed LM. Joint OPT wins because the gate fires on the trigger N-gram (a deterministic surface property of the question) and the inserted row biases the next-token distribution toward the gold first token; retrieval can miss the relevant turn (visible in RAG\_TOP1’s 0.101) and, when it hits, the in-context evidence still competes with the base LM’s distributional priors.

**Scaling at the optimal config.** We repeated the same LOCOMO setup with all four trained Mini-Engrams at our ablation-optimal recipe (Section ??); the Engram lead widens monotonically with dense size (Figure ??).

**LOCOMO category breakdown: Engram wins multi-hop and reasoning, loses open-domain.** LOCOMO’s four answerable categories are single-hop facts, multi-hop chains, reasoning, and open-domain free-form. The previous tables reported only single-hop. Running all 3 dense scales across all 4 categories produces a more nuanced picture (Table ??):

**The Engram win is category-dependent, not uniform.** Multi-hop and reasoning categories favour Engram by 49–178% because the per-fact (question, first-answer-token) OPT loss directly encodes the question→answer mapping, which retrieval struggles to chain across evidence sentences. **Open-domain flips the story:** questions like “What did the charity race raise awareness for?” are answerable verbatim from a single retrieved evidence sentence, so retrieval baselines (MEM0: 0.247–

LOCOMO single-hop: token-F1 over-credits Engram; semantic judge flips the ranking




Figure 7: Locomo single-hop, full 10 conversations, scaling with Mini-Engram dense size. **(a) token-F1**: User-as-Engram Joint OPT (blue) beats every retrieval baseline at every scale. **(b) LLM-judge accuracy** (Qwen2.5-14B-Instruct): the picture flips—retrieval baselines (MEMO\_LIKE, MEMMACHINE) beat Joint OPT by 0.05–0.10 because token-F1 over-credits Engram for the correct first answer token despite a noisy continuation.




Figure 8: Locomo token-F1 vs. LLM-judge accuracy at every system  $\times$  every dense size. Retrieval baselines (red/orange) sit near the  $y=x$  line: their token-F1 and LLM-judge scores roughly agree. User-as-Engram OPT and Joint OPT (blue/cyan) sit systematically *above* the line, indicating that token-F1 over-credits Engram. The asymmetry is the metric artifact documented in Section ??.

0.341, MEMMACHINE: 0.209–0.338) decisively beat Engram (0.122–0.159) because Engram only conditions a single next-token bias and cannot reconstruct an unseen evidence sentence from that signal.

This category breakdown reframes the contribution: User-as-Engram is a *specialised* memory mechanism that excels at chain-of-fact recall but is not a universal replacement for retrieval. A production system should route by question type or compose Engram (for trained facts) with retrieval (for free-form synthesis from session context).

**LLM-judge correction: retrieval wins on token-F1’s blind spot.** Token-F1 rewards partial word overlap. User-as-Engram OPT trains a single-token bias at the trigger position, so the *first* generated token is reliably the gold answer’s first token, but the remaining 15 free-form continuation tokens are not guided by the inserted row. To check whether token-F1 over-credits this first-token bias, we

Table 12: LOCOMO category  $\times$  model token-F1. Engram Joint OPT vs. the best retrieval baseline (MEMO\_LIKE or MEMMACHINE\_LIKE) per cell. The Engram win is category-dependent.

| category           | answering LM | MEMO         | MEMMACH      | Engram J-OPT | lead vs. best retr. |
|--------------------|--------------|--------------|--------------|--------------|---------------------|
| single-hop         | d12 v2       | 0.131        | 0.116        | <b>0.169</b> | +29%                |
|                    | d12@1280     | 0.160        | 0.152        | <b>0.176</b> | +10%                |
|                    | d20          | 0.161        | 0.169        | <b>0.233</b> | +38%                |
| multi-hop          | d12 v2       | 0.092        | 0.076        | <b>0.243</b> | <b>+165%</b>        |
|                    | d12@1280     | 0.115        | 0.104        | <b>0.253</b> | <b>+120%</b>        |
|                    | d20          | 0.102        | 0.108        | <b>0.299</b> | <b>+178%</b>        |
| reasoning          | d12 v2       | 0.117        | 0.113        | <b>0.183</b> | +56%                |
|                    | d12@1280     | 0.123        | 0.136        | <b>0.203</b> | +49%                |
|                    | d20          | 0.135        | 0.168        | <b>0.266</b> | +58%                |
| <b>open-domain</b> | d12 v2       | <b>0.247</b> | 0.209        | 0.122        | − <b>51%</b>        |
|                    | d12@1280     | <b>0.341</b> | 0.327        | 0.146        | − <b>57%</b>        |
|                    | d20          | 0.314        | <b>0.338</b> | 0.159        | − <b>53%</b>        |

Table 13: LLM-as-judge accuracy (Qwen2.5-14B-Instruct) on the full 10-conversation LOCOMO single-hop set ( $\sim 800$  QAs per model). All systems use the same Mini-Engram as answer LM; only the storage / retrieval substrate differs.

| answering LM                | NO_MEM | MEMO         | MEMMACH      | OPT   | Joint-OPT | best retr. lead  |
|-----------------------------|--------|--------------|--------------|-------|-----------|------------------|
| Mini-Engram-d8 v2           | 0.005  | <b>0.116</b> | 0.106        | 0.028 | 0.044     | +0.07 over J-OPT |
| Mini-Engram-d12 v2          | 0.010  | 0.156        | <b>0.158</b> | 0.041 | 0.059     | +0.10 over J-OPT |
| Mini-Engram-d12@1280        | 0.025  | <b>0.169</b> | 0.158        | 0.079 | 0.095     | +0.07 over J-OPT |
| <b>Mini-Engram-d20@1536</b> | 0.030  | <b>0.190</b> | 0.177        | 0.110 | 0.140     | +0.05 over J-OPT |

re-scored every prediction with **Qwen2.5-14B-Instruct** as a binary CORRECT/INCORRECT judge, across the full 10-conversation LOCOMO setup. The picture flips (Table ??):

**Under semantic LLM-judge, retrieval beats Engram on LOCOMO single-hop at every dense size.** The lead shrinks with scale (+0.07 at d8 to +0.05 at d20) but does not reverse. Token-F1 over-credited Engram by 0.05–0.10 because it rewards the correct first answer token without penalising a noisy continuation (Figure ??; Engram points sit above the  $y=x$  line, retrieval points sit on it). *This is a correction to the headline of Table ??.* Engram still wins paraphrase queries (96.9% top-1 vs MEMMACHINE’s 75.0%, Table ??; unchanged by the judge correction since paraphrase scores single-token recall) and storage (88 KB / 100 facts, metric-independent).

**Multi-token OPT closes the judge gap.** The gap is driven by per-fact OPT training only the first answer token; the remaining 15 continuation tokens are unguided. Replacing  $-\log p(y_0 \mid \text{trigger})$  with the autoregressive sum  $-\sum_t \log p(y_t \mid \text{trigger}, y_{<t})$  pushes the residual stream toward states that produce the *full* answer (Table ??).

**At d20 under matched multi-token pipelines, J-OPT-multi reaches token-F1 0.310 vs. 0.197 best retrieval (+57%) and LLM-judge 0.225 vs. 0.174 (+29%).** Engram-winning kicks in between d12 v2 and d12@1280; below that, dense capacity is too small for the continuation to fall into place after the Engram-anchored prefix.

**Limitations of this evaluation.** The judge is Qwen2.5-14B-Instruct under a binary CORRECT/INCORRECT prompt (stricter than LOCOMO’s 0–5 official judge), so direct comparison to published LOCOMO leaderboards requires re-judging with the canonical pipeline. The LLM-judge coverage here is single-hop only; temporal/adversarial questions are not scored, and the multi-hop limitation of Section ?? applies to chained queries without suffix-N-gram overlap.

Table 14: Single-hop LOCOMO (full 10 conversations,  $\sim 800$  QAs per model). Multi-token Engram Joint-OPT closes the judge gap. **All systems are scored under one matched pipeline**: the retrieval baselines and Engram J-OPT-multi both use the same multi-token answer-generation pipeline (greedy 16-token continuation, Qwen2.5-14B LLM-judge). J-OPT-first is reported under the first-token pipeline for reference. The previous version of this table reported retrieval judge scores under the first-token pipeline (a more favourable comparison for retrieval); we now report both pipelines but use the matched pipeline as the headline.

| answering LM                | token-F1 (multi-token pipeline) |         |             |                    | LLM-judge accuracy (matched, multi-token) |         |             |                     |
|-----------------------------|---------------------------------|---------|-------------|--------------------|-------------------------------------------|---------|-------------|---------------------|
|                             | MEM0                            | MEMMACH | J-OPT first | <b>J-OPT multi</b> | MEM0                                      | MEMMACH | J-OPT first | <b>J-OPT multi</b>  |
| Mini-Engram-d8 v2           | 0.116                           | 0.103   | 0.134       | <b>0.159</b>       | 0.081                                     | 0.076   | 0.044       | 0.060               |
| Mini-Engram-d12 v2          | 0.142                           | 0.134   | 0.169       | <b>0.249</b>       | 0.128                                     | 0.145   | 0.059       | 0.105               |
| Mini-Engram-d12@1280        | 0.181                           | 0.168   | 0.176       | <b>0.229</b>       | 0.125                                     | 0.117   | 0.095       | <b>0.159</b> (+27%) |
| <b>Mini-Engram-d20@1536</b> | 0.197                           | 0.196   | 0.233       | <b>0.310</b>       | 0.174                                     | 0.168   | 0.140       | <b>0.225</b> (+29%) |

Bold marks cells where J-OPT-multi beats the best retrieval baseline. The cross-over to Engram-winning occurs between d12 v2 and d12@1280. Under the prior (unmatched) reporting where retrieval was scored under its better first-token pipeline, retrieval judge scores were MEM0/MEMMACH = 0.116/0.106, 0.156/0.158, 0.169/0.158, 0.190/0.177 across the four scales; that conservative comparison shrinks Engram’s d20 lead from +29% (matched) to +18% (unmatched).

Table 15: Capacity levels used in the ablation. Total Engram params =  $4 \cdot v \cdot e$  (2 N-gram orders  $\times$  2 layers).

| label  | $v$ slots | $e$ embed | total Engram |
|--------|-----------|-----------|--------------|
| tiny   | 5 K       | 64        | 1.28 M       |
| small  | 20 K      | 128       | 10.24 M      |
| medium | 50 K      | 128       | 25.6 M       |
| large  | 50 K      | 256       | 51.2 M       |
| xlarge | 100 K     | 256       | 102.4 M      |

## 6.8 Engram capacity ablation

How does Engram-table size interact with token budget at a fixed dense backbone? We probe this as a  $5 \times 3 \times 2$  matrix: five capacity levels ( $v$  slots per N-gram,  $e$  total embed-dim per N-gram), three token budgets, two dense sizes (Figure ??).

**Three regularities.** (i) *Capacity has an optimum.* Both endpoints under-perform: *tiny* can’t address enough distinct facts, *xlarge* under-trains each slot. (ii) *The optimum is the same Engram size at both dense scales:* large ( $50\text{ K} \times 256 = 51.2\text{ M}$  params) wins at the highest token budget for both d8 and d12. **The optimum scales in tokens, not capacity.** (iii) *Below the catch-up token budget, smaller Engrams win.* At d8 / 0.5 B tokens, *small* (10 M) beats *large* (51.2 M) on ins-OPT 1.00 vs. 0.62—the larger table is under-trained per slot. The cross-over happens around 1 B tokens for d8 and 1.32 B tokens for d12.

**Practical rule.** For an Engram pretrained at Karpathy 12 t/p of scaling-params, use the *large* ( $50\text{ K} \times 256$ ) Engram table from d12@768 onward. For d8-class models with  $\leq 1\text{ B}$  training tokens, use *small* ( $20\text{ K} \times 128$ ) instead. Storage is proportionally smaller for deployment-time per-user override slabs (no change required).

## 6.9 Fact-count scaling: per-fact independent OPT to 1000 facts

Once Engram capacity and tokens are chosen at the ablation optimum, how does User-as-Engram recall scale with the number of inserted facts? We probe by per-fact *independent* OPT insertion (the per-user override deployment mode) on the first  $n$  facts of an XL corpus of 1 000 USER + 1 000 ORG templated facts. ICL@1000 is the in-context ceiling at the hardest end. Figure ?? shows the full curves.

**Per-fact recall is approximately flat in  $n$  up to 1000 facts.** The best d8 and d12@1280 cells hold  $\geq 0.98$  top-1 from  $n=100$  to  $n=1000$ . The crucial caveat is that this is *independent* per-fact OPT: each fact gets its own private hash rows, with no within-table interference between facts—the per-user

## Engram capacity $\times$ tokens response surface




Figure 9: Engram capacity  $\times$  token-budget response surface, at Mini-Engram-d8 v2 (left) and d12 v2 (right). Top row: E1 USER OPT top-1 (100-fact recall). Bottom row: LOCOMO Joint OPT F1. The same large (50 K  $\times$  256) Engram size wins at the highest token budget for *both* dense scales; *tiny* is under-capacity, *xlarge* is over-provisioned. The optimum scales in tokens, not capacity.

Table 16: Engram capacity  $\times$  token budget at Mini-Engram-d8 v2 (137 M dense, 42 M scaling-params). Cells report ins-OPT 16-fact top-1 / E1 USER OPT 100-fact top-1.

| capacity | 0.5 B tok   | 1.0 B tok   | 2.0 B tok          |
|----------|-------------|-------------|--------------------|
| tiny     | 0.44 / 0.62 | 0.56 / 0.57 | 0.69 / 0.69        |
| small    | 1.00 / 0.95 | 1.00 / 0.98 | 1.00 / 1.00        |
| medium   | 0.88 / 0.84 | 1.00 / 0.97 | 1.00 / 1.00        |
| large    | 0.62 / 0.84 | 1.00 / 1.00 | <b>1.00 / 1.00</b> |
| xlarge   | 0.75 / 0.67 | 0.81 / 0.80 | 0.81 / 0.79        |

override deployment mode of Section ?? . In that regime, we observe no recall ceiling up to 1,000 facts on a 625 M-parameter base LM. Joint OPT into a shared table (Section ??) is the regime where the density ceiling appears (68%  $\rightarrow$  35% top-1 from 100 to 1,000 facts at d12@768).

### 6.10 Dense-size scaling at optimal config

Pulling the four trained Mini-Engrams into one row each (Figure ??):

**Three observations.** (i) **Val bpb scales smoothly with dense  $\times$  tokens:** 0.91  $\rightarrow$  0.83  $\rightarrow$  0.77  $\rightarrow$  0.73. (ii) **Insertion-OPT and E1 saturate** from d12@768 onward. The 16-fact ins-OPT probe reaches 1.00 at every dense size with  $\geq$  1.32 B tokens; E1 USER/ORG OPT reaches 1.00 at d12@1280@3.34 B, dipping slightly to 0.97 at d20@1536@7.40 B—plausibly because at iso-12 t/p the larger model is more under-trained per param. (iii) **LOCOMO Joint OPT first-token F1 scales monotonically with dense size** from 0.134 (178 M) to 0.233 (1.22 B), and the multi-token variant



Table 17: Engram capacity  $\times$  token budget at Mini-Engram-d12@768 (339 M dense, 110 M scaling-params). Cells report ins-OPT 16-fact top-1 / E1 USER OPT 100-fact top-1.

| capacity | 0.5 B tok   | 1.32 B tok  | 2.5 B tok          |
|----------|-------------|-------------|--------------------|
| tiny     | 0.31 / 0.25 | 0.44 / 0.40 | 0.62 / 0.57        |
| small    | 0.81 / 0.78 | 0.81 / 0.92 | 1.00 / 0.99        |
| medium   | 0.88 / 0.74 | 0.81 / 0.81 | 0.94 / 0.97        |
| large    | 0.94 / 0.93 | 1.00 / 0.98 | <b>1.00 / 1.00</b> |
| xlarge   | 1.00 / 0.96 | 0.94 / 0.95 | 0.94 / 0.97        |

Per-fact independent OPT recall is flat in  $n$  up to 1000 facts




Figure 10: Per-fact independent OPT recall vs. fact count  $n$ . Left: USER OPT top-1. Right: ORG OPT top-1. The d12@1280 and d20@1536 optimal cells (green, purple) hold near-1.00 recall from  $n = 100$  to  $n = 1000$ . No ceiling visible in the per-fact independent-OPT regime.

climbs from 0.159 to 0.310 (Table ??). Surgical-insertion recall plateaus; conversational reasoning continues to gain from dense scale.

**Reading the gap.** Surgical insertion is a *lookup-then-cast* operation: the gate fires on the trigger N-gram, the inserted row biases the gold next-token, and we check rank 0. Once the base LM is fluent enough to make the gate selective ( $\geq$  d12@768), recall saturates. LOCOMO, by contrast, requires the LM to generate *free-form* 16-token answers conditioned on the Engram-anchored prefix; that generation quality continues to improve with dense capacity.

### 6.11 Multi-hop reasoning over inserted facts

Per the discussion of LoRA’s multi-hop failure (Section ??), we directly probe Engram’s behaviour on chained facts. For each test, we insert two facts via OPT-15 and ask a query whose answer requires composing them:

```
Fact 1: My doctor's name is Dr.      -> Patel
Fact 2: Dr. Patel works at           -> Globex
Query:  My doctor works at           -> ? (expected: Globex)
```

We probe Engram’s multi-hop behaviour on a corpus of 63 chained-fact items spanning 13 schema families, with each item labelled by whether the query’s suffix N-gram *overlaps* Fact 2’s trigger (Table ??). The split is near-balanced (32 overlap, 31 no-overlap), which lets us isolate the surface-trigger mechanism from true compositional chaining.

**The 8-pair 75% headline was an overlap-heavy sample.** On the balanced 63-pair corpus, multi-hop top-1 drops to **52%** with a much tighter Wilson 95% CI [40.2%, 64.4%]. Decomposing by overlap status reveals a near-bimodal mechanism: **surface-overlap items reach 91%** (CI [75.8%, 96.8%]), while **no-overlap items collapse to 13%** (CI [5.1%, 28.9%]) and even there 2 of the 4 “successes” are first-fact-wins coincidences where the trained Fact 1 row’s strong logit also covers the query (e.g., “*My favorite spice’s source country is*” returning “India” for every spice in the family). True

Table 18: E1 USER OPT top-1 / top-5 vs. fact count  $n$ . Per-fact independent OPT-15 insertion. Last column: ICL ceiling at  $n = 1000$ .

| model                                  | params        | $n=100$     | $n=200$     | $n=500$     | $n=1000$           | ICL@1000    |
|----------------------------------------|---------------|-------------|-------------|-------------|--------------------|-------------|
| Mini-Engram-d8 v1                      | 137 M         | 0.81 / 0.91 | 0.82 / 0.92 | 0.81 / 0.93 | 0.84 / 0.95        | 0.98 / 1.00 |
| Mini-Engram-d8 large/2 B (best d8)     | 178 M         | 1.00 / 1.00 | 0.99 / 1.00 | 0.99 / 1.00 | 0.99 / 1.00        | 0.99 / 1.00 |
| Mini-Engram-d12 v1                     | 339 M         | 0.86 / 0.96 | 0.88 / 0.96 | 0.87 / 0.97 | 0.88 / 0.97        | 0.96 / 1.00 |
| Mini-Engram-d12 v2 (12 t/p)            | 339 M         | 0.86 / 0.98 | 0.87 / 0.97 | 0.86 / 0.97 | 0.87 / 0.97        | 1.00 / 1.00 |
| Mini-Engram-d12 large/2.5 B (best d12) | 339 M         | 0.96 / 1.00 | 0.95 / 1.00 | 0.95 / 0.99 | 0.96 / 0.99        | 0.98 / 1.00 |
| Mini-Engram-d12@1280 opt.              | 625 M         | 0.98 / 1.00 | 0.97 / 1.00 | 0.98 / 1.00 | 0.98 / 1.00        | 1.00 / 1.00 |
| <b>Mini-Engram-d20@1536 opt.</b>       | <b>1.22 B</b> | 0.90 / 0.99 | 0.92 / 0.99 | 0.91 / 0.99 | <b>0.93 / 0.99</b> | 0.99 / 1.00 |

Dense-size scaling at the ablation-optimal recipe




Figure 11: Dense-size scaling at the ablation-optimal recipe. Left: Locomo Joint OPT first-token token-F1 climbs from 0.134 (d8 v2) to 0.233 (d20@1536); MEMMACHINE\_LIKE retrieval baseline climbs more slowly because its quality is bounded by sentence-encoder retrieval, not the LM. Right: val\_bpb (grey) drops monotonically with dense+tokens, while E1 USER OPT single-fact recall (blue) saturates at 1.00 from d12@1280 onward (small dip at d20 at iso-12 t/p).

compositional chaining without surface overlap is 10%, consistent with the architectural prediction that the gate is a surface-N-gram hash and does not compose across triggers. Per-fact direct recall sanity is 99.2%, so the multi-hop gap is purely the chaining-vs-lookup distinction, not a Fact-1 insertion failure.

**Conclusion.** User-as-Engram solves chained queries *when* the query’s suffix matches an inserted trigger (91%), and effectively fails on true compositional chains (13%, dropping to 10% once first-fact-wins coincidences are filtered). This is the architectural limit shared with LoRA-as-memory: a surface-N-gram gate cannot compose across triggers by construction. The recite-then-reason recipe (Section ??) is the natural mitigation; multi-fact-in-the-loss pretraining (Section ??) is the structural fix. The original 75% pilot pre-dated this decomposition; it is preserved as a row in Table ?? for traceability.

## 6.12 LoRA rank ablation at 100 facts

For completeness we also ablate the multi-fact LoRA rank at fixed training budget (2000 steps, lr 5e-4, all  $Q/K/V$  projections):

Every rank  $\geq 8$  saturates at the recall ceiling with a 2000-step budget. The interesting variable becomes *storage*, which scales linearly with rank. The point: at 100 facts/user, all rank- $\geq 8$  LoRAs hit the recall ceiling, so the LoRA-vs-Engram trade-off collapses to recall-vs-storage: **Engram Joint OPT is 20 $\times$  smaller than the smallest LoRA at the ceiling (88 KB vs. 1.8 MB rank-8), in exchange for a  $\sim 35$ -pt top-1 gap (65% vs. 100%; top-5 gap is 2 pt: 98% vs. 100%).** That gap closes by  $n=1\,000$  where LoRA rank-64 also slips to 38% top-1 (Section ??).

Table 19: Mini-Engram scaling. Four rows are at the ablation-optimal recipe (large Engram, Karpathy 12 t/p  $\times$  scaling-params): d8 v2, d12 v2, **d12@1280 opt.**, and **d20@1536 opt.** The two large/2B and large/2.5B rows are overtrained reference points ( $4\times$  and  $\sim 2\times$  the 12 t/p budget respectively), included to show per-fact recall saturates at 1.00 once tokens exceed Chinchilla-optimal. All trained from scratch on ClimbMix-400B, bf16, single Blackwell GPU. “Total” includes the 51.2 M Engram-table params; “scaling” is the dense transformer body (the figure that drives the 12 t/p schedule). LOCOMO J is Joint OPT first-token token-F1 on the full 10-conversation single-hop set, matching Table ??.

| model                            | total         | scaling      | tokens        | val bpb      | ins-OPT     | E1 USER     | E1 ORG      | LOCOMO J     |
|----------------------------------|---------------|--------------|---------------|--------------|-------------|-------------|-------------|--------------|
| Mini-Engram-d8 v2                | 178 M         | 42 M         | 0.50 B        | 0.912        | 0.62        | 0.81        | 0.78        | 0.134        |
| Mini-Engram-d8 large/2B          | 178 M         | 42 M         | 2.00 B        | —            | 1.00        | 1.00        | 1.00        | —            |
| Mini-Engram-d12 v2               | 339 M         | 110 M        | 1.32 B        | 0.827        | 1.00        | 0.98        | 0.93        | 0.169        |
| Mini-Engram-d12 large/2.5B       | 339 M         | 110 M        | 2.50 B        | —            | 1.00        | 1.00        | 1.00        | —            |
| <b>Mini-Engram-d12@1280 opt.</b> | <b>625 M</b>  | <b>278 M</b> | <b>3.34 B</b> | <b>0.770</b> | <b>1.00</b> | <b>1.00</b> | <b>1.00</b> | <b>0.176</b> |
| <b>Mini-Engram-d20@1536 opt.</b> | <b>1.22 B</b> | <b>617 M</b> | <b>7.40 B</b> | <b>0.730</b> | <b>1.00</b> | 0.97        | 0.97        | <b>0.233</b> |

Table 20: Multi-hop reasoning over Engram-inserted facts ( $n=63$  chained pairs on Mini-Engram-d20). The Engram rows insert two facts per item via OPT-15 and ask the chained query. Per-fact direct recall sanity-checks that both facts are insertable ( $125/126 = 99.2\%$ ); top-1 and top-5 differ only marginally (33 vs 35), confirming the original observation that the gate either fires on the second-hop trigger and returns the gold token, or it does not fire and the gold is far from the top. The **surface-overlap vs no-overlap split** measures the mechanism directly. The 8-pair RAG sub-block below is from the original pilot for context only; the load-bearing RAG comparison is Section ??’s KB-scale study.

| subset (Mini-Engram-d20)                                                              | per-fact direct (sanity) | multi-hop top-1        | multi-hop top-5 | Wilson 95% CI on top-1 |
|---------------------------------------------------------------------------------------|--------------------------|------------------------|-----------------|------------------------|
| surface-overlap items (query suffix = Fact 2 trigger)                                 | 64 / 64 (100%)           | <b>29 / 32 (90.6%)</b> | 29 / 32 (90.6%) | [75.8%, 96.8%]         |
| no-overlap items (true chaining required)                                             | 61 / 62 (98.4%)          | <b>4 / 31 (12.9%)</b>  | 6 / 31 (19.4%)  | [5.1%, 28.9%]          |
| <b>combined (<math>n=63</math>)</b>                                                   | <b>125 / 126 (99.2%)</b> | <b>33 / 63 (52.4%)</b> | 35 / 63 (55.6%) | [40.2%, 64.4%]         |
| <i>Original 8-pair pilot (no-overlap items: “My doctor works at”; “My pet eats”):</i> |                          |                        |                 |                        |
| Mini-Engram-d20 (8-pair pilot)                                                        | 16 / 16 (100%)           | 6 / 8 (75.0%)          | 6 / 8 (75.0%)   | [40.9%, 93.0%]         |
| <i>RAG over the original 8 pair-facts, Mini-Engram-d20 backbone (context):</i>        |                          |                        |                 |                        |
| Mini-Engram-d20 + RAG top-1                                                           | 0 / 8 (0%)               | 4 / 8 (50.0%)          | 6 / 8 (75.0%)   | —                      |
| Mini-Engram-d20 + RAG top-2                                                           | 8 / 8 (100%)             | 8 / 8 (100%)           | 8 / 8 (100%)    | —                      |
| Mini-Engram-d20 + RAG top-2 + shared LoRA $r=16$                                      | 8 / 8 (100%)             | 7 / 8 (87.5%)          | 8 / 8 (100%)    | —                      |

### 6.13 Inserted-fact recall in free continuation

Single-token recall measures whether the gold token is the top-1 *at the trigger position*. In conversational use the model continues for many tokens; we test whether the inserted fact surfaces in the first 8 generated tokens (greedy decoding) on 30 facts after Joint OPT (1500 steps).

**Dense scaling sharpens generation behaviour.** At d12 v1 the gold token is the immediate top-1 only 23% of the time but *anywhere in the next 8 tokens* 53% of the time—the model “approaches” the fact across a few tokens. From d12@1280 onward, the gold either lands at position 0 or never appears in the window: the bigger model commits earlier rather than drifting. We caution that  $n=30$  is small (Wilson 95% CI on 27/30 is [74%, 97%]; on 18/30 is [42%, 75%]), and the result is consistent with the bigger model simply missing the fact when its first guess is wrong. Either way, the 90% first-token rate at d20@1536 is the practical recall in conversational use.

### 6.14 Paraphrase generalisation

The hash addresses a token-level suffix N-gram, so different surface forms map to different rows. The full comparison against retrieval baselines under paraphrase queries is in Table ??; multi-trigger insertion reaches 96.9% top-1 vs. MEMMACHINE’s 75.0%, and the lead is scale-invariant (Section ??, “Scale invariance” paragraph). Per-paraphrase rank data is in Appendix ??.

Table 21: Multi-fact LoRA at 100 facts on Mini-Engram-d12@1280 (625 M), 2000 steps. Storage assumes fp32.

| LoRA rank          | 8     | 32     | 64     | 128    |
|--------------------|-------|--------|--------|--------|
| top-1              | 1.00  | 1.00   | 1.00   | 0.99   |
| top-5              | 1.00  | 1.00   | 1.00   | 1.00   |
| parameters         | 442 K | 1.77 M | 3.54 M | 7.08 M |
| storage (MB, fp32) | 1.8   | 7.1    | 14.2   | 28.3   |

Table 22: Free-continuation recall: does the inserted fact surface in the first 8 generated tokens after Joint-OPT?

| answering LM                              | first-token rate   | anywhere-in-8 rate |
|-------------------------------------------|--------------------|--------------------|
| Mini-Engram-d12 v1 (339 M)                | 7/30 = 23%         | 16/30 = 53%        |
| Mini-Engram-d12@1280 opt. (625 M)         | 18/30 = 60%        | 18/30 = 60%        |
| <b>Mini-Engram-d20@1536 opt. (1.22 B)</b> | <b>27/30 = 90%</b> | <b>27/30 = 90%</b> |

## 7 Multi-tenant serving system

We implement `EngramServer`, a  $\sim 50$ -line wrapper around an Engram-pretrained model that supports per-user (and per-org) override maps with sub-millisecond apply/restore. The full per-request latency CDF is in Appendix Figure ??.

**Architecture.** HBM holds the frozen base + global Engram tables. DRAM holds the override store keyed by `(scope_id)`  $\rightarrow$  `OverrideMap`. Per request: resolve user/org  $\rightarrow$  apply overrides  $\rightarrow$  forward  $\rightarrow$  restore. There is no router, no fused LoRA kernel, no graph rewrite.

**Empirical results** (30 users  $\times$  50 facts  $\times$  600 requests on d12, Table ??):

The system scales linearly in tenants. At the optimal d12@1280, 100 users  $\times$  100 facts gives 60% top-1 / 97% top-5 under the production serving protocol with 1000-step Joint OPT (Table ??); recall is bounded by within-user fact density, not by tenant count, and top-5 holds at 97% even at the largest tested density. Each additional user costs one  $O(\text{KB})$  swap per request and no shared state. The  $70\times$  override-apply speedup from d12 v1 to d12@1280 (at 30 u/50 f) comes from loading the larger model fully on-device and avoiding a PCIe round-trip.

**Cross-user privacy.** The *observed* “leak” rate in cross-user-probe mode (deliberately serving user U with user V’s question) depends on the protocol. Under the proper override apply/restore (production serving), V’s overrides are not in the table when V is not the active user, so the leak comes only from gold-value coincidence between users (1.3–3.9% at 20–30 users; 4.4% at 100 u  $\times$  100 f under Joint-OPT, dominated by users sharing the same favourite spice or surname). Under the all-coresident protocol where every user’s overrides are written into one shared table simultaneously, the apparent leak rises to 46% at 100 u  $\times$  100 f because identical gold-value surface forms collide across users in the same address space. Cross-user information flow under the serving protocol is *zero by construction* regardless of pool size.

**LoRA-serving comparison.** S-LoRA (?) and Punica (?) serve per-user LoRAs via custom CUDA kernels with per-batch-row routing. `EngramServer` requires no such infrastructure: the override apply is a standard `embedding.weight[rows] = ...` operation. Production batched-multi-user serving needs one extra gather op per Engram layer; we do not benchmark that kernel here.

## 8 Layered architecture: content + meta-skill

So far: per-user LoRA contaminates (Section ??) and per-user Engram preserves direct recall but lacks in-context reasoning skill (23% indirect\_any alone at  $n=20$ , Section ??). Can the two be *composed* without inheriting LoRA’s damage? Our thesis: personal memory decomposes into *content*

Figure 5: EngramServer overhead is < 10% of request latency.




Figure 12: Multi-tenant serving on Mini-Engram-d12 (30 users  $\times$  50 facts  $\times$  600 requests). **(a)** Override apply latency distribution: median 2.2 ms. **(b)** Per-request latency breakdown: apply (2.2 ms) + forward (18.7 ms) + restore (2.3 ms) = 23.2 ms median. Override overhead is <10% of request latency.

Table 23: EngramServer evaluation across four deployment scales, single Blackwell GPU. d12@1280 columns are the ablation-optimal Mini-Engram (625 M). <sup>†</sup>The 100 u/100 f registration time was measured on a GPU shared with another workload; idle-GPU registration extrapolates to  $\sim 8$ –10 s/user from the 30 u/50 f baseline. Throughput and end-to-end latency are omitted from this column for the same reason; override-apply timing (CPU-side) and recall (which is GPU-contention-invariant) are reported.

| metric                        | d8 (20 u/30 f)                 | d12 v1 (30 u/50 f) | d12@1280 (30 u/50 f)  | d12@1280 (100 u/100 f) |
|-------------------------------|--------------------------------|--------------------|-----------------------|------------------------|
| registration time / user      | 15 s                           | 20 s               | <b>7.0 s</b>          | 14.8 s <sup>†</sup>    |
| storage / user                | 20 KB                          | 59 KB              | 59 KB                 | 89 KB                  |
| throughput (1-tok recall)     | 42.8 req/s                     | 47.9 req/s         | <b>232 req/s</b>      | — <sup>†</sup>         |
| latency p50 / p99             | 23.8 / 33.1 ms                 | 23.2 / 27.8 ms     | <b>4.2 / 4.4 ms</b>   | — <sup>†</sup>         |
| override apply p50 / p99      | 0.4 / 3.3 ms                   | 2.2 / 2.3 ms       | <b>0.03 / 0.04 ms</b> | 2.25 / 2.27 ms         |
| own-fact recall top-1 / top-5 | 83% / 100%                     | 76% / 98%          | <b>76% / 99%</b>      | <b>60% / 97%</b>       |
| cross-user privacy            | <b>0% leak by construction</b> |                    |                       |                        |

(per-user, trigger-keyable, low-cost) and *meta-skill* (cross-user reasoning patterns, learnable from held-out users); matching each to its right substrate strictly dominates every single-substrate baseline. Two hypotheses make this concrete:

- **H1 (naïve combination)**: per-user LoRA + per-user Engram (condition D) inherits LoRA’s indirect-reasoning damage; Engram’s local edit does not rescue the composition.
- **H2 (layered architecture)**: one *shared* LoRA (cross-user meta-skill) + per-user Engram (local content) decomposes memory correctly. The shared LoRA’s contamination is amortised across the population, and the per-user Engram adds none on top ( $\Delta\text{bpb}(F) = \Delta\text{bpb}(E)$  by Section ??’s locality property).

## 8.1 Design

Train a single shared LoRA on cross-user in-context-reasoning data: for each held-out training user (here, u020–u029), render each indirect QA as a completion-format sample “Facts: <fact1>. <fact2>. ... Q: <q> A: <gold>”, using only the facts in `required_fact_keys`. This teaches the LoRA the *pattern* “given facts in context, do the reasoning,” without memorising any specific user’s content. At inference, attach the shared LoRA + the test user’s Engram-row override and ask the indirect question without facts in the prompt: the Engram provides content (via trigger N-gram lookup), the shared LoRA provides the skill.

Table 24: Layered architecture: 10 conditions  $\times$  20 test users on Mini-Engram-d20 (canonical seed S0; all numbers in this table from a single LoRA-init seed so RAG rows G–J are matched to parametric rows A–F). All training is independent (per-user LoRA, per-user Engram, and shared LoRA each trained alone). RAG rows (G–J) retrieve facts at inference with all-MiniLM-L6-v2 and prepend them in the same Facts: . . . template as the shared-LoRA training data. Rows G–I have no parametric edit, so  $\Delta\text{bpb} = 0$  by construction; the cost they incur is the context tokens column instead. “worse” = fraction of users with indirect\_any below the no-edit baseline. Appendix ?? reports two additional seeds (S1, S2) of conditions A–F: 3-seed mean F indirect\_any is 41% (range 35–45%), 3-seed mean B indirect\_any is 7% (range 6–9%).

| code                                                                                                   | retrieval / shared-LoRA<br>+ per-user edit | direct<br>top-1 | direct<br>top-5 | indirect<br>top-1 | indirect<br>any                  | $\Delta\text{bpb}$<br>(worse/20) | ctx<br>toks |
|--------------------------------------------------------------------------------------------------------|--------------------------------------------|-----------------|-----------------|-------------------|----------------------------------|----------------------------------|-------------|
| A                                                                                                      | —                                          | 29%             | 38%             | 14%               | 19%                              | 0.000 (0/20)                     | 0           |
| B                                                                                                      | per-user LoRA $r=64$                       | 99%             | 99%             | 38%               | <b>6%</b> $\downarrow\downarrow$ | <b>+1.784 (17/20)</b>            | 0           |
| C                                                                                                      | per-user Engram J-OPT                      | 100%            | 100%            | 14%               | 23%                              | +0.00005 (0/20)                  | 0           |
| D                                                                                                      | per-user LoRA + Engram                     | 100%            | 100%            | 37%               | <b>8%</b> $\downarrow\downarrow$ | +1.819 (18/20)                   | 0           |
| E                                                                                                      | shared LoRA $r=16$ only                    | 54%             | 76%             | 62%               | 44% $\uparrow\uparrow$           | +0.386 (0/20)                    | 0           |
| F                                                                                                      | <b>Engram + shared LoRA</b>                | <b>100%</b>     | <b>100%</b>     | <b>61%</b>        | <b>44%</b>                       | <b>+0.386 (0/20)</b>             | <b>0</b>    |
| <i>RAG conditions (this work): retrieve user’s facts at inference, no parametric edit on the base.</i> |                                            |                 |                 |                   |                                  |                                  |             |
| G                                                                                                      | RAG top-1                                  | 81%             | 90%             | 16%               | 33%                              | 0.000 (0/20)                     | 27          |
| G′                                                                                                     | oracle top-1 (gold fact only)              | 84%             | 94%             | 14%               | 26%                              | 0.000 (0/20)                     | 28          |
| H                                                                                                      | RAG top-3                                  | 77%             | 89%             | 16%               | 39%                              | 0.000 (0/20)                     | 44          |
| I                                                                                                      | RAG all ( $\approx$ MARKDOWN_ALL)          | 76%             | 85%             | 15%               | 35%                              | 0.000 (0/20)                     | 302         |
| J                                                                                                      | <b>RAG top-3 + shared LoRA</b>             | 77%             | 91%             | <b>80%</b>        | <b>54%</b>                       | +0.386 (0/20)                    | 44          |

## 8.2 Six-condition head-to-head

We evaluate every per-user combination on  $n=20$  test users ( $n=10$  training users held out for the shared LoRA), measuring direct top-1/top-5 recall on the user’s facts (completion-format prompts), indirect top-1 and indirect-any on completion-format indirect probes (“Q: <q>\nA:”, greedy 16-token continuation), and val\_bpb on a held-out 262 K-token ClimbMix shard (Table ??).

## 8.3 Results: layered design Pareto-dominates

**F vs B (the headline contrast).** The layered design matches per-user LoRA on direct recall (F 100% vs B 99%) while delivering **7.4 $\times$  higher indirect\_any (44% vs 6%)**, **4.6 $\times$  less contamination ( $\Delta\text{bpb} +0.39$  vs  $+1.78$ )**, and **0/20 users worse than the no-edit base** on indirect (vs 17/20 for per-user LoRA). Per-user storage drops from 14.2 MB to the Engram’s 88 KB because the cross-user LoRA is amortised. Paired-bootstrap 95% CI on the per-user  $F - B$  indirect\_any delta (3 seeds pooled  $\rightarrow n=60$ , 2000 resamples):  $[+31.2, +36.5]$  pp (mean  $+33.8$  pp).

**H1 (naive combination) is CONFIRMED.** Adding Engram on top of per-user LoRA (condition D) does *not* rescue indirect reasoning: 18/20 users worse than base on S0; 50/60 = 83% pooled across seeds, comparable to LoRA-alone’s 49/60. The contamination dominates regardless of the local edit added on top.

**H2 (layered architecture) is CONFIRMED on all three sub-claims.** (i) *Locality survives layering*:  $\Delta\text{bpb}(F) = \Delta\text{bpb}(E)$  in every seed ( $+0.379/+0.386/+0.388$  at S2/S0/S1)—adding Engram contributes zero contamination on top of the shared LoRA. (ii) *Direct recall is preserved*: F is 100% in every seed, matching C—Engram does the lookup unaffected by the shared LoRA. (iii) *Meta-skill enables indirect reasoning*: F’s indirect\_any (S0 44%, 3-seed mean 41%) is  $\sim 2\times$  both the per-user Engram baseline (C 23%) and the no-edit baseline (A 19%). Neither C alone (no in-context skill) nor E alone (54% direct, never saw the user’s facts) suffices; the combination is what works.

## 8.4 Does RAG close the indirect-reasoning gap?

The natural rebuttal to F is: *just retrieve the facts and put them in the base LM’s context*. Conditions G–J in Table ?? test this with all-MiniLM-L6-v2 retrieval at  $k \in \{1, 3, \text{all}\}$ , prepending the retrieved facts in the same Facts: . . . Q: . . . A: template used for shared-LoRA training. Three findings:




Figure 13: Indirect-reasoning accuracy (`indirect_any`) vs. avg context tokens per query on the 20-user split. F (layered) and E (shared LoRA only) anchor the 0-context frontier; J (RAG top-3 + shared LoRA) lifts accuracy to 54% at 44 context tokens; Qwen-3B + RAG (Section ??) reaches 55–57% at 82–390 tokens on a  $\sim 2\times$  larger backbone. Naive RAG on the same Mini-Engram-d20 base (G/H/I) sits *below* F: putting facts in context is not free meta-skill.

- **Naive RAG on a base LM is below F on indirect.** Conditions G–I yield 33–39% `indirect_any`, below F’s 44%. An Engram-pretrained base never learned to consult an in-context fact list, so it ignores the prefix and falls back to parametric priors. Direct recall under RAG also collapses to 76–84% (vs F’s 100%) for the same reason.
- **Oracle retrieval doesn’t help.** G’ uses ground-truth `required_fact_keys` (no retriever noise) and still gets 26% `indirect_any`—worse than RAG top-1’s 33%. The bottleneck is not retrieval quality; the base LM can’t use a single fact for indirect reasoning without the meta-skill.
- **Shared LoRA + RAG (J) is the strongest indirect setting at non-zero context cost.** J reaches 54% `indirect_any` and 80% `indirect_top1`, exceeding F (44%/61%) by +10 pp / +19 pp—but at 44 context tokens per query instead of zero.

This refutes the “just use RAG” rebuttal on the same Engram-pretrained base and identifies a clean Pareto choice (Figure ??): zero context budget  $\rightarrow$  F; unrestricted context  $\rightarrow$  J.

## 8.5 RAG with an instruction-tuned LM head (Qwen-3B)

The Mini-Engram-d20 base above is a base LM and does not follow ICL well, which is what makes conditions G–I weak. The fairest test of “RAG alone is enough” plugs retrieved facts into a real instruction-tuned LM. We use Qwen2.5-3B-Instruct ( $\sim 3$  B parameters;  $\sim 2.5\times$  Mini-Engram-d20), the same encoder (all-MiniLM-L6-v2), and the same 20 users  $\times$  20 indirect probes. The chat-template prompt gives Qwen the retrieved facts as a system-bulleted list and asks it to answer the user’s question (Appendix ??).

Qwen-3B + RAG-all hits 57% `indirect_any` (its retrieval ceiling); RAG top-1 lands at 55%. Our **J configuration on the 1.22 B Mini-Engram-d20 hits 54%**—within 3 pp of a strictly larger instruction-tuned model at  $2\times$  less context. Our **F (no retrieval, 0 context tokens) is 12 pp below Qwen-3B + RAG-all**, on a base that was never instruction-tuned. On `indirect_top1`, J’s 80% exceeds every Qwen+RAG variant.

Table 25: Indirect reasoning on the same 20-user/20-probe split as Table ??, but with Qwen2.5-3B-Instruct as the answering LM and retrieval over the user’s facts. The published F (Mini-Engram-d20, layered) line is reproduced for direct comparison. “ret\_acc” is the fraction of probes where the retrieved set covers all gold required\_fact\_keys.

| method                                                 | indirect top-1 | indirect any | ret_acc | ctx toks |
|--------------------------------------------------------|----------------|--------------|---------|----------|
| Qwen-3B, no context                                    | 16%            | 16%          | 0%      | 67       |
| Qwen-3B + RAG top-1                                    | 50%            | 55%          | 22%     | 82       |
| Qwen-3B + RAG top-3                                    | 50%            | 53%          | 62%     | 102      |
| Qwen-3B + RAG all (34 facts)                           | <b>55%</b>     | <b>57%</b>   | 100%    | 390      |
| Qwen-3B + oracle top-1                                 | 42%            | 46%          | 37%     | 83       |
| <i>Same backbone family (Mini-Engram-d20, 1.22 B):</i> |                |              |         |          |
| F: layered (Engram + shared LoRA), this work           | 61%            | 44%          | —       | <b>0</b> |
| J: RAG top-3 + shared LoRA, this work                  | <b>80%</b>     | 54%          | 62%     | 44       |




Figure 14: Indirect-reasoning accuracy and retrieval recall vs. KB size ( $n=20$  users, 20 indirect probes each; KB = test user’s 34 facts + distractors sampled from the 30-user schema-family pool, log  $x$ -axis,  $N \in \{34, 100, 200, 300, 500, 1000\}$ ). (a) F (layered, horizontal solid line at 44%) is invariant to KB size; the retrieval-based conditions degrade monotonically. **Naïve RAG and Qwen-3B + RAG fall below F at KB  $\geq 100$ ; at KB=1000 the gap is 14 pp.** J (RAG top-3 + shared LoRA) is the most robust to KB growth (54%  $\rightarrow$  47%) because the shared meta-skill compensates when retrieval misses, though its lead over F shrinks from 10 pp to 3 pp across the sweep. (b) Retrieval recall (fraction of probes where the retrieved set covers all gold required\_fact\_keys) drops by 7 $\times$  for top-3 (62%  $\rightarrow$  9%) between KB=34 and KB=1000, which is the mechanistic driver of (a).

**Interpretation.** Qwen-3B + RAG is the strongest realistic “user-as-RAG” baseline and tops out at 55–57% indirect\_any. F (layered) hits 44% *without retrieval*; J (RAG + shared LoRA) closes the gap when context is free. This is the Pareto picture in Figure ??.

**Wall-clock per query** (same RTX 6000): F  $\sim 310$  ms (16-token greedy, no retrieval), J  $\sim 440$  ms (with RAG top-3), Qwen-3B + RAG top-3  $\sim 92$  ms (smaller dense-only). Latency is comparable; the deployment-scale gap is in per-user storage (88 KB Engram vs.  $\sim 130$  stored fact strings, Figure ??).

## 8.6 RAG-vs-Engram trade-offs sharpen as the KB grows

Tables ?? and ?? used each user’s native 34 facts. Production fact counts are typically larger ( $10^2$ – $10^3$ : chat history, calendar, preferences, contacts). We augment each test user’s 34 facts with distractors sampled from the 30-user schema-family pool (u000–u029 minus the active user; 1008 facts) to reach  $N \in \{34, 100, 200, 300, 500, 1000\}$ . Indirect probes are unchanged; the retriever now discriminates the gold fact among  $N$  candidates. F is unaffected: the per-user Engram table holds only the test user’s facts, independent of population size.



Table 26: RAG indirect\_any (%) vs. KB size on the same 20-user split. “ret%” = fraction of probes with all required\_fact\_keys retrieved at the indicated  $k$ . Mini-Engram-d20 F (layered) is unchanged across KB sizes (44% indirect\_any, 0 context tokens).

| method                                                 | KB=34 |           | KB=100 |           | KB=200 |           | KB=300 |           | KB=500 |           | KB=1000 |           |
|--------------------------------------------------------|-------|-----------|--------|-----------|--------|-----------|--------|-----------|--------|-----------|---------|-----------|
|                                                        | ret%  | any%      | ret%   | any%      | ret%   | any%      | ret%   | any%      | ret%   | any%      | ret%    | any%      |
| Mini-Engram-d20 + RAG top-1 (G)                        | 22    | 33        | 13     | 31        | 9      | 30        | 7      | 30        | 7      | 30        | 6       | 29        |
| Mini-Engram-d20 + RAG top-3 (H)                        | 62    | 39        | 39     | 35        | 28     | 32        | 20     | 30        | 12     | 30        | 9       | 30        |
| Mini-Engram-d20 + RAG top-3 + shared LoRA (J)          | 62    | <b>54</b> | 39     | <b>52</b> | 28     | <b>51</b> | 20     | <b>51</b> | 12     | <b>48</b> | 9       | <b>47</b> |
| Qwen-3B + RAG top-1                                    | 22    | 55        | 13     | 41        | 9      | 35        | 7      | 33        | 7      | 32        | 6       | 32        |
| Qwen-3B + RAG top-3                                    | 62    | 52        | 39     | 45        | 28     | 38        | 20     | 37        | 12     | 32        | 9       | 30        |
| <b>F: layered (Engram + shared LoRA), no retrieval</b> | —     | <b>44</b> | —      | <b>44</b> | —      | <b>44</b> | —      | <b>44</b> | —      | <b>44</b> | —       | <b>44</b> |

Table 27: Shared-LoRA rank ablation on Mini-Engram-d20 ( $n=20$  test users, all 20 users of the headline split). All cells use independent training (shared LoRA trained once on u020–u029; per-user Engram J-OPT trained per test user). The layered design (F) preserves direct recall at every rank while indirect performance and  $\Delta$ bpb peak together at  $r=16$ .

| shared rank              | direct top-1 (F) | indirect top-1 (F) | indirect any (F) | $\Delta$ bpb (F) |
|--------------------------|------------------|--------------------|------------------|------------------|
| $r=4$                    | 100%             | 52%                | 28%              | +0.302           |
| <b><math>r=16</math></b> | <b>100%</b>      | <b>62%</b>         | <b>44%</b>       | <b>+0.386</b>    |
| $r=64$                   | 100%             | 36%                | 35%              | +1.027           |

Retrieval degrades fast: top-3 recall on all-MiniLM-L6-v2 drops  $7\times$  ( $62\% \rightarrow 9\%$ ) purely from more candidates in the same embedding space; top-1 drops  $22\% \rightarrow 6\%$ . Naive RAG falls 14–15 pp below F at  $N \geq 100$ . Qwen-3B + RAG, the strongest realistic baseline, starts above F at KB=34 (52%) but crosses below at KB=100 and ends 14 pp below F at KB=1000. J (RAG + shared LoRA) is the most KB-robust RAG configuration, but its lead over F shrinks from 10 pp at KB=34 to 3 pp at KB=1000—and J pays 42 context tokens plus a shared-LoRA forward pass that F does not. F holds 44% across the sweep.

**Substrate ranking shifts with deployment scale.** At KB=34, Qwen-3B + RAG is the accuracy ceiling and F is the zero-context option. At KB  $\geq 200$ , F decisively beats both naive RAG (29–32%) and Qwen-3B + RAG (30–38%); the “RAG wins at context cost” trade-off holds only at small KB.

**Where each side’s failure comes from.** The two limits are mechanistically distinct. Engram’s density ceiling (Section ??) comes from forward-pass interference among co-active rows inside one user’s table; it is bounded by per-user fact count. RAG’s retrieval ceiling comes from nearest-neighbour confusion across an ever-growing candidate pool; it is bounded by whatever pool the retriever sees (per-user, per-tenant, or per-corpus). For  $10^2$ – $10^3$  facts/user, Engram density is the binding constraint on the parametric side; for  $10^4$ + facts in a shared corpus, retrieval is the binding constraint on the RAG side. The two costs scale on different axes and cross around  $N \approx 100$ .

## 8.7 Rank ablation: $r=16$ is the sweet spot

We sweep the shared LoRA rank  $r \in \{4, 16, 64\}$  on the full 20-user split (Table ??) and observe a clear, non-monotonic curve. Smaller ranks under-fit the meta-skill; larger ranks both contaminate more (more parameters, more global perturbation) *and* learn the meta-skill worse, presumably because the LoRA’s capacity exceeds what the 510-sample training corpus can profitably constrain.

The  $r=64$  row is the surprise: a  $4\times$  larger shared LoRA gives *worse* indirect performance (35% vs 44% at  $r=16$ ) at  $2.7\times$  more contamination. We recommend  $r=16$ ; future work could revisit this once the shared training corpus is scaled beyond 510 samples.

## 8.8 Cross-schema generalisation

The headline F vs B numbers train the shared LoRA on held-out users in the same synthetic schema as the test users. We test whether the meta-skill transfers *across* schemas, the condition that killed Stage A’s recite-then-reason recipe in Section ?? (within-schema 0.41, cross-schema 0.046—a 90% collapse). We train the rank-16 shared LoRA on the original first-person personal-facts schema




Figure 15: Cost-quality Pareto for the layered architecture (Section ??). Same Mini-Engram-d20 base,  $n=20$  test users, six substrate combinations. **F (LAYERED) is on the Pareto frontier**: matches the highest indirect-reasoning accuracy (44% indirect-any) at low storage (88 KB/user) and low contamination ( $\Delta\text{bpb} +0.39$ ). Per-user LoRA (B) is Pareto-dominated by F on every axis. Naive combination (B+C, the “add Engram on top of LoRA” baseline) is also dominated.

Table 28: Cross-schema head-to-head on Mini-Engram-d20 ( $n=20$  test users in the medical schema; shared LoRA trained on  $n=10$  users in the personal schema). The layered F vs per-user-LoRA B contrast preserves its qualitative shape: F matches direct recall (100%) at  $7.6\times$  higher indirect\_any (31% vs 4%) and  $3.3\times$  less contamination (+0.39 vs +1.28). Compare with Section ??’s Stage A recite-then-reason cross-schema collapse ( $0.41 \rightarrow 0.046$ ).

| code     | condition                   | direct top-1 | direct top-5 | indirect top-1 | indirect any | $\Delta\text{bpb}$ |
|----------|-----------------------------|--------------|--------------|----------------|--------------|--------------------|
| A        | no edit                     | 21%          | 35%          | 19%            | 29%          | 0.000              |
| B        | per-user LoRA $r=64$        | 100%         | 100%         | 50%            | <b>4%</b> ↓↓ | <b>+1.283</b>      |
| C        | per-user Engram J-OPT       | 100%         | 100%         | 19%            | 29%          | +0.00001           |
| D        | per-user LoRA + Engram      | 100%         | 100%         | 46%            | <b>4%</b> ↓↓ | +1.477             |
| E        | shared LoRA $r=16$ only     | 43%          | 48%          | 72%            | 31% ↑↑       | +0.386             |
| <b>F</b> | <b>Engram + shared LoRA</b> | <b>100%</b>  | <b>100%</b>  | <b>72%</b>     | <b>31%</b>   | <b>+0.386</b>      |

(u020–u029, “My name is...”, “I was born in...”) and evaluate the full six-condition head-to-head on a held-out medical schema (m000–m019, third-person patient framing: “The patient’s name is...”, “MRN: ...”; 29 facts, 20 indirect probes per user) (Table ??).

**The layered claim survives cross-schema, with a realistic decay.** F’s indirect\_any degrades from 44% (within-schema, Table ??) to 31% (cross-schema, this table)—a 30% relative drop. Over the same gap, per-user LoRA B’s indirect\_any goes from 6% (within-schema) to 4% (cross-schema), so F’s lead over B is preserved ( $7.4\times$  within-schema,  $7.6\times$  cross-schema). The architectural locality is invariant by construction:  $\Delta\text{bpb}(F) = \Delta\text{bpb}(E) = +0.386$  in both regimes, identical to four decimal places, because the per-user Engram still fires only on its trigger N-grams. *Architectural decomposition transfers across schemas; recipe-level fixes (Stage A) did not.*

Table 29: Layered head-to-head on a Q/A-format-adapted Mini-Engram-d20 ( $n=20$  test users u000–u019). “ $\Delta\text{bpb}$ ” is relative to the SFT’d base’s val\_bpb (4.42). All edits drop val\_bpb because the SFT moved the base away from raw ClimbMix; the conserved quantity is  $|\Delta\text{bpb}(F) - \Delta\text{bpb}(E)|$ , which measures the contamination added by Engram on top of shared LoRA and is  $1 \times 10^{-4}$  on this base—essentially identical to the  $5 \times 10^{-5}$  noise floor on the original d20 base (Section ??).

| code     | condition                   | direct top-1 | direct top-5 | indirect top-1 | indirect any | $\Delta\text{bpb}$ |
|----------|-----------------------------|--------------|--------------|----------------|--------------|--------------------|
| A        | no edit                     | 12%          | 16%          | 1%             | 1%           | 0.0000             |
| B        | per-user LoRA $r=64$        | 100%         | 100%         | 24%            | 6%           | −0.441             |
| C        | per-user Engram J-OPT       | 97%          | 98%          | 1%             | 1%           | +0.0022            |
| D        | per-user LoRA + Engram      | 100%         | 100%         | 24%            | 3%           | −0.517             |
| E        | shared LoRA $r=16$ only     | 51%          | 69%          | 36%            | 16%          | −1.4268            |
| <b>F</b> | <b>Engram + shared LoRA</b> | <b>100%</b>  | <b>100%</b>  | <b>37%</b>     | <b>16%</b>   | <b>−1.4267</b>     |

### 8.9 Layered F on a Q/A-format-adapted Engram base

The headline F vs B numbers are on a base LM (Mini-Engram-d20). Section ??’s cross-base table shows that per-user LoRA’s *user-visible* damage shrinks dramatically on instruction-tuned bases (85% of users worse on a base LM, 0–20% on four instruction-tuned bases). Does F still dominate B when B itself becomes a stronger baseline?

We SFT Mini-Engram-d20 on  $\sim 2000$  single-token (Q, A) pairs from the same fact templates the layered experiment uses (1500 steps,  $\text{lr } 1 \times 10^{-4}$ , post-embedding params trained, Engram tables frozen). The SFT is intentionally narrow—no general-text corpus mixed in—so it tests Q/A format adaptation, not full chat-tuning. val\_bpb on raw ClimbMix rises from 0.73 to 4.42, indicating the model has been pushed toward the Q/A distribution at the cost of general-text LM ability. We then train a fresh rank-16 shared LoRA on this SFT’d base (shared\_lora\_d20\_sft/r16, 2000 steps on u020–u029) and rerun the six-condition head-to-head (Table ??).

**F still Pareto-dominates B; the lead shrinks as B strengthens.** The qualitative picture survives Q/A adaptation: F matches B on direct (100% vs 100%), beats B on indirect ( $1.6\times$  on top-1,  $2.8\times$  on indirect\_any), and the architectural locality is preserved ( $|\Delta\text{bpb}(F) - \Delta\text{bpb}(E)| = 1 \times 10^{-4}$ , on the same order as the  $5 \times 10^{-5}$  original-base noise floor). The cross-base prediction reproduces: per-user LoRA’s indirect\_top1 jumps from 1.3% (A) to 23.5% (B) on this adapted base, matching Section ??’s “instruction-tuned bases absorb the LoRA perturbation” pattern. F’s lead over B on indirect\_any compresses from  $7.4\times$  (base LM) to  $2.8\times$  (SFT’d), but F is still the dominant substrate.

**Caveats.** The SFT here is narrow Q/A adaptation, not full chat-tuning: mixing general-text in the SFT corpus or starting from a true instruction-tuned base (the cross-base table’s Qwen/Llama/Mistral setup, but with an Engram-pretrained backbone) is the natural follow-up. The aggressive SFT also collapses A and C’s indirect to floor (1%), which inflates F’s relative lead on top of A but not the F-vs-B contrast that matters for the substrate-selection conclusion.

### 8.10 Storage and amortisation

Per user, the layered design costs the Engram’s 88 KB; the shared LoRA is one global 11.8 MB (rank-16) artifact amortised across the population. For 1 M users that’s  $100 \text{ GB} + 12 \text{ MB} \approx 100 \text{ GB}$ , indistinguishable from Engram alone, vs. 14.2 TB for per-user LoRA at the same recall ceiling. Figure ?? plots all six substrate combinations on a (storage, indirect-reasoning) axis.

### 8.11 Limitations of the layered measurement

(i) Cross-schema generalisation was the load-bearing open question; Section ?? shows F’s lead survives a personal  $\rightarrow$  medical schema shift with a 30% relative degradation ( $44\% \rightarrow 31\%$  indirect\_any), versus the 90% collapse on Stage A’s recite-then-reason recipe in Section ?. We tested one schema pair; broader domain shifts (cross-language, technical jargon) are open. (ii) The headline base is a base LM (Mini-Engram-d20); Section ?? shows F still Pareto-dominates B on a Q/A-format-adapted

Mini-Engram-d20 ( $2.8\times$  on indirect\_any), though the lead compresses from  $7.4\times$  as B becomes a stronger baseline. A full chat-tuned Engram-pretrained backbone (the cross-base table’s setup) is the natural next test. (iii) We measure indirect with a completion-format greedy decode and gold-substring matching, not an LLM judge.

## 9 Discussion and limitations

### 9.1 When to use what

Substrate selection turns on three axes: *per-user fact count*, *whether queries need indirect reasoning*, and *whether context budget or per-user storage is the binding constraint*. Table ?? gives the recommended substrate per regime.

Table 30: Substrate-selection cheat sheet. Numbers cite this work unless marked external. “invariant” = the metric does not degrade with KB / population size.

| regime                                    | recommended substrate            | per-user cost             | headline metric                | key citation |
|-------------------------------------------|----------------------------------|---------------------------|--------------------------------|--------------|
| <10 facts/user, direct recall only        | ICL / strong-retriever RAG       | 0 KB persistent           | near-perfect recall            | Sec. ??      |
| 10–1000 facts/user, indirect matters      | layered F (Engram + shared LoRA) | <b>88 KB</b> , 0 ctx toks | 44% indirect (invariant in KB) | Sec. ??–??   |
| 10–1000 facts/user, ctx budget free       | J (RAG top-3 + shared LoRA)      | 88 KB + 44 ctx toks       | 54% indirect (KB=34)           | Sec. ??      |
| enterprise, 100% direct top-1, <10K users | per-user LoRA (S-LoRA serving)   | 14.2 MB                   | 100% direct, indirect tax      | Sec. ??      |
| long-form / open-domain / shared corpus   | RAG                              | 0 KB persistent           | wins LOCOMO open-domain +53%   | Tab. ??      |

The load-bearing row is the second: layered F is the right substrate for the dominant regime (10–1000 facts/user, indirect reasoning matters, many users), and decisively beats RAG with a  $\sim 2.5\times$  larger instruction-tuned LM at  $N \geq 100$ . Any substrate hitting 100% direct top-1 incurs an indirect-reasoning tax; the question is whether that tax is structural (global LoRA), constructive (local Engram, with the trigger N-gram as contract), or retrieval-bounded (RAG, with the encoder discriminating gold among  $N$  candidates).

### 9.2 Open limitations

**Shared multi-hop reasoning gap.** Both LoRA and Engram are surface-trigger keyed; neither composes facts across triggers (“if my doctor is Patel and Patel works at Globex, what is my doctor’s employer?”). User-as-Engram inherits this gap from User-as-LoRA. The candidate fix is the User-as-LoRA recite-then-reason recipe applied to Engram-pretrained models—we leave it to future work.

**Within-user density ceiling.** Joint OPT at 1000 facts gets 35% top-1. The slot space is highly sparse (<1% occupancy), so the constraint is forward-pass interference among co-active rows, not hash collisions. Per-user scoped tables (load only the relevant rows for the current query, via a small classifier) is a candidate future fix; another is the multi-fact-in-the-loss recipe change in Section ??.

**Engram pretraining required.** Our method assumes an Engram-pretrained base. We trained four ourselves at 178 M, 339 M, 625 M, and 1.22 B total parameters; production deployment depends on either DeepSeek releasing weights or pretraining your own ( $\sim 10$  h on a single Blackwell GPU at our 625 M scale; production-scale Engrams require correspondingly larger compute budgets).

### 9.3 Future work: training-recipe changes

We trained the d8 v2, d12 v2, d12@1280, and d20@1536 Mini-Engrams ourselves at the ablation-optimal recipe (Section ??); LOCOMO Joint OPT first-token F1 scales monotonically with dense size to 0.233 at 1.22 B (and multi-token to 0.310). Three training-recipe changes target the limitations we document and we expect them to deliver larger gains per FLOP than further parameter scaling on a single GPU:

- *Multi-fact insertion in the loss.* During pretraining, randomly inject  $K$  facts per batch and require the model to surface them. This trains the gate to handle multi-fact density natively, addressing the within-user joint-density ceiling visible in Table ?? (35% top-1 at 1000 facts).

- *Random per-batch hash salts.* Salt 50% of batches with a random per-batch user salt during pretraining. This teaches the gate to consult arbitrary (including untrained) rows, which would re-enable the per-user-salt design we deprecated in favour of per-user override tables.
- *Recite-then-reason traces.* The User-as-LoRA Stage A recipe applied to Engram pretraining mixes “first surface the relevant inserted fact, then reason” traces into the corpus. This is the most direct attack on the shared multi-hop gap.

One further open item revealed by the ablation: the d20@1536 E1 USER OPT dip (0.97 vs. d12@1280’s 1.00) at iso-12 t/p suggests bigger models need higher t/p to fully saturate insertion recall. Note that the Joint-OPT density ceiling *does not* relax with dense scale at fixed Engram capacity (Table ??: top-1 at 1 000 facts actually degrades from 0.35 at d12@768 to 0.28 at d12@1280 and d20@1536), so breaking it likely requires the *multi-fact-in-the-loss* recipe change above rather than further dense scaling.

**Teacher-trace supervision (negative result).** We tested whether retraining the shared LoRA on a teacher-distilled trace corpus (observation / factual-statement / QA /  $\langle$ think $\rangle$ -reasoning, 245 Qwen3-8B traces) would improve F. Two variants both reduced contamination ( $\Delta$ bpb +0.386  $\rightarrow$  +0.233 / +0.196, 40–49% less) but *also* reduced indirect-any accuracy (44.5%  $\rightarrow$  38.5%/28.0%; paired-bootstrap 95% CIs strictly negative). ?’s trace recipe is co-designed with an agent-format eval protocol that scores committed fields after an optional  $\langle$ think $\rangle$  block; our layered evaluation scores top-tokens of a direct Q: . . . A: completion. Composability would require either retargeting the trace supervision to a top-token objective or extending our eval to a  $\langle$ think $\rangle$  slot—parametric-memory recipes are not eval-protocol-agnostic.

**Layered-architecture follow-ups.** Two remaining directions (cross-schema and a Q/A-format-adapted base addressed in Sections ??–??): (i) scale the shared-LoRA training corpus beyond 510 samples to test whether the rank-64 over-parameterisation (Table ??) reverses; (ii) jointly train the per-user Engram with the shared LoRA (currently independent) as a probe for additional synergy. A full chat-tuned Engram base (starting from a chat corpus rather than Q/A pairs) remains open.

## 10 Conclusion

**Personal memory is two problems, not one: content and meta-skill.** Matching each to its right substrate Pareto-dominates every per-user single-substrate baseline we tested. Three findings carry the paper:

**(1) The architectural contamination gap is large and seed-stable:**  $\sim 34,000\times$  on Mini-Engram-d20 ( $\Delta$ bpb +1.78 for per-user LoRA vs +0.00005 for Engram-row insertion). User-visible damage is base-dependent (85% worse on the base LM; 0–20% on four instruction-tuned bases); the architectural gap is not, and it motivates a local-edit substrate.

**(2) The layered design Pareto-dominates:** F (shared LoRA + per-user Engram) matches per-user LoRA’s direct recall (100% vs 99%) at  $7.4\times$  higher indirect reasoning (44% vs 6%),  $4.6\times$  less contamination, and 0/60 (vs 49/60 pooled) users hurt on indirect. The lead survives a personal  $\rightarrow$  medical schema shift ( $7.6\times$  over B cross-schema, identical locality), unlike the recite-then-reason recipe which collapsed 90% under the same shift. The naive composition (LoRA + Engram) inherits LoRA’s damage (50/60 worse).

**(3) Substrate ranking is decided by deployment scale.** At KB=34, Qwen-3B + RAG slightly beats F (52% vs 44%); at KB  $\geq 100$ , F is invariant while Qwen-3B + RAG drops to 30% (14 pp below F), because the per-user Engram table does not grow with population size. On LOCOMO, per-user Engram wins reasoning (+58%) and multi-hop (+178%, via surface overlap) and—with multi-token OPT under matched pipelines—single-hop on LLM-judge (+29%) at 1.22 B; it loses open-domain (–53%), where retrieving an unseen evidence sentence is required.

Three recommendations follow. **(i) Decompose by substrate:** content  $\rightarrow$  per-user Engram override (88 KB/user); meta-skill  $\rightarrow$  one shared LoRA amortised over users; open-domain synthesis  $\rightarrow$  RAG. **(ii) Match the training objective to the deployment metric:** multi-token OPT for free-form answers, first-token OPT when the inserted span suffices. **(iii) Substrate selection, not gradient training, is the bottleneck of personal memory:** gate-vs-store separation, layered composition, and per-user

override tables together deliver zero-leak personal memory at  $\sim 100$  GB total storage for 1 M users and 232 req/s on one GPU.

## References

- X. Cheng, W. Zeng, D. Dai, Q. Chen, B. Wang, Z. Xie, K. Huang, X. Yu, Z. Hao, Y. Li, H. Zhang, H. Zhang, D. Zhao, and W. Liang. Conditional memory via scalable lookup: A new axis of sparsity for large language models. arXiv:2601.07372, January 2026.
- A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. Attention is all you need. NeurIPS 2017.
- T. Brown et al. Language models are few-shot learners. NeurIPS 2020.
- A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever. Language models are unsupervised multitask learners (GPT-2). OpenAI technical report, 2019.
- S. Min, X. Lyu, A. Holtzman, M. Artetxe, M. Lewis, H. Hajishirzi, and L. Zettlemoyer. Rethinking the role of demonstrations: What makes in-context learning work? EMNLP 2022.
- P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W. Yih, T. Rocktäschel, S. Riedel, and D. Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. NeurIPS 2020.
- Y. Gao et al. Retrieval-augmented generation for large language models: A survey. arXiv:2312.10997, 2023.
- N. Shazeer, A. Mirhoseini, K. Maziarz, A. Davis, Q. Le, G. Hinton, and J. Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. ICLR 2017.
- D. Dai et al. DeepSeekMoE: Towards ultimate expert specialization in mixture-of-experts language models. arXiv:2401.06066, 2024.
- S. Mangrulkar et al. PEFT: State-of-the-art parameter-efficient fine-tuning methods. <https://github.com/huggingface/peft>, 2022.
- K. Jordan et al. Muon: An optimiser for hidden layers in neural networks. GitHub / blog post, 2024.
- C. Huang et al. LoRAHub: Efficient cross-task generalization via dynamic LoRA composition. arXiv:2307.13269, 2023.
- D. Wu et al. LongMemEval: Benchmarking chat assistants on long-term memory. arXiv 2024.
- A. Maharana et al. LOCOMO: Evaluating very long-term conversational memory of LLM agents. ACL 2024.
- M. Tavakoli, A. Salemi, C. Ye, M. Abdalla, H. Zamani, and J. R. Mitchell. Beyond a million tokens: Benchmarking and enhancing long-term memory in LLMs (BEAM). arXiv:2510.27246, 2025.
- B. Jiang, Y. Yuan, M. Shen, Z. Hao, Z. Xu, Z. Chen, Z. Liu, A. R. Vijjini, J. He, H. Yu, R. Poovendran, G. Wornell, L. Ungar, D. Roth, S. Chen, and C. J. Taylor. PersonaMem-v2: Towards personalized intelligence via learning implicit user personas and agentic memory. arXiv:2512.06688, 2025.
- A. Karpathy. nanochat: an experimental training harness for LLMs. <https://github.com/karpathy/nanochat>, 2026.
- E. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen. LoRA: Low-rank adaptation of large language models. ICLR 2022.
- N. Houlsby et al. Parameter-efficient transfer learning for NLP. ICML 2019.
- W. Su et al. Parametric retrieval-augmented generation (PRAG). arXiv:2501.15915, 2025.
- Z. Tan et al. DyPRAG: Dynamic parametric RAG. arXiv:2505.19386, 2025.

- J. Chen, H. Zhang, L. Pang, Y. Tong, H. Zhou, Y. Zhan, W. Lin, and Z. Zheng. Privacy-preserving reasoning with knowledge-distilled parametric retrieval-augmented generation (DistilledPRAG). arXiv:2509.01088, 2025.
- Z. Tan, Q. Liu, and M. Jiang. Democratizing large language models via personalized parameter-efficient fine-tuning (OPPU). EMNLP 2024 (arXiv:2402.04401).
- Y. Zhuang et al. HYDRA: Per-user adapters for personalised LLMs. arXiv 2024.
- M. Bini, O. Bohdal, U. Michieli, Z. Akata, M. Ozay, and T. Ceritli. MemLoRA: Distilling expert adapters for on-device memory systems. arXiv:2512.04763, 2025.
- R. Charakorn, E. Cetin, Y. Tang, and R. T. Lange. Text-to-LoRA: Instant transformer adaption. arXiv:2506.06105, 2025.
- Z. Tan et al. PER-PCS: Per-user post-hoc LoRA composition. arXiv 2024.
- Y. Sheng, S. Cao, D. Li, et al. S-LoRA: Serving thousands of concurrent LoRA adapters. arXiv:2311.03285, 2024.
- L. Chen, Z. Ye, Y. Wu, et al. Punica: Multi-tenant LoRA serving. MLSys 2024.
- G. Lample, A. Sablayrolles, M. Ranzato, L. Denoyer, and H. Jégou. Large memory layers with product keys. NeurIPS 2019.
- P. He. PEER: Mixture of one million experts. arXiv 2024.
- V. Berges, B. Oğuz, D. Haziza, W. Yih, L. Zettlemoyer, and G. Ghosh. Memory layers at scale. ICML 2025.
- Z. Huang, Q. Min, H. Huang, D. Zhu, Y. Zeng, R. Guo, and X. Zhou. Ultra-sparse memory network (Ultra-Mem). arXiv:2411.12364, 2024 (ICLR 2025).
- J. Huang et al. OverEncoding: hashed N-gram embeddings via averaging. 2025.
- L. Yu et al. SCONE: scalable contextual N-gram embeddings. 2025.
- A. Pagnoni, R. Pasunuru, P. Rodriguez, et al. BLT: byte latent transformer with hashed N-gram embeddings. arXiv:2412.09871, 2025.
- A. Liu et al. SuperBPE: word-level BPE for compositional patterns. 2025.
- CivitAI Community. LoRA stacking patterns for Stable Diffusion. <https://civitai.com/>, 2024.
- K. Meng, D. Bau, A. Andonian, and Y. Belinkov. Locating and editing factual associations in GPT (ROME). NeurIPS 2022.
- K. Meng et al. MEMIT: Mass-editing memory in a transformer. ICLR 2023.
- R. Cohen et al. Evaluating the ripple effects of knowledge editing in language models (MQuAKE). 2023.
- R. Cohen et al. RippleEdits: A benchmark for ripple effects of model editing. 2024.
- K. Meng et al. CounterFact: a counterfactual editing benchmark. 2022.
- D. Wu, J.-C. Gu, K.-W. Chang, and N. Peng. Self-routing RAG: Binding selective retrieval with knowledge verbalization. arXiv:2504.01018, 2025.
- J. Hewitt et al. Introspective reasoning traces for parametric knowledge surfacing. arXiv:2602.22193, 2026.
- S. Gekhman et al. Thinking to Recall: chain-of-thought pre-recitation for multi-hop parametric recall. arXiv:2603.09906, 2026.
- Z. Sun et al. Recitation-augmented language models. ICLR 2023.

Table 31: Per-user direct vs. indirect recall, POLAR LoRA Stage A on Qwen2.5-3B-Instruct (10 users, 50 facts each). Sample of users 0–1; the original  $n=10$  pilot mean was  $\Delta=-0.008$  (6/10 users worse than base). At the scaled  $n=30$  replication (Table ?? in the main paper) the mean inverts to  $\Delta=+0.088$  with 6/30 users worse, refuting the original headline.

| uid                                     | direct (with adapter) | indirect (with adapter) | indirect (base only) | train sec |
|-----------------------------------------|-----------------------|-------------------------|----------------------|-----------|
| u000                                    | 1.000                 | 0.133                   | 0.200                | 102       |
| u001                                    | 1.000                 | 0.071                   | 0.214                | 98        |
| ... (8 more users with similar pattern) |                       |                         |                      |           |
| <b>mean (n=10)</b>                      | <b>1.000</b>          | <b>0.150</b>            | <b>0.170</b>         | 100       |

Table 32: Stage A baselines (mean over 10 users, 30 indirect Qs each).

| condition                           | direct       | indirect |
|-------------------------------------|--------------|----------|
| no insertion (base only)            | 0.000        | 0.170    |
| POLAR adapter                       | <b>1.000</b> | 0.150    |
| POLAR adapter + explicit CoT prompt | —            | 0.296    |
| ICL upper bound (facts in context)  | 0.515        | 0.304    |

B. Li. Hippo: Autonomous Parametric Learning from Agent Experience. Pine AI technical report, 2026. Concurrent work; companion paper to this one. Identifies two failure modes of canonical NTP-based parametric memory (agent-format surfacing penalty + Allen-Zhu Part 3.2 manipulation collapse) and introduces a fourth NTP training format — teacher-distilled  $\langle \text{think} \rangle$  reasoning traces — that fixes both, with multi-seed paired-bootstrap CIs on a per-user LoRA substrate. We test their recipe on our layered architecture in Section ?? (negative result).

## A Appendix: Detailed User-as-LoRA results

This appendix expands Section ?? with the full empirical record that motivates the present paper. All results are on Qwen2.5-3B-Instruct with a rank-64 per-user LoRA (POLAR-class), 10 synthetic users, and the trace-mix Stage A recipe where indicated. The reader should treat this as the empirical foundation for our claim that LoRA-as-memory is reasoning-negative.

### A.1 Stage A pilot: per-user breakdown

In 5 of 10 users the indirect recall with the adapter is *strictly worse* than the no-adapter base (Figure ??b in main paper). The direct recall is saturated at 1.000 for every user.

### A.2 Stage A baselines

The CoT prompt + adapter condition reaches the ICL ceiling, showing the facts are *accessible* through the adapter—just not spontaneously invoked. This rules out “the adapter doesn’t contain the answer” as the cause of the indirect-recall gap.

### A.3 Trace-mix Stage A: within-schema vs. cross-schema

Within-schema indirect recovers to 0.41 (Wilson 95% CI [0.30, 0.54]), matching the ICL-with-CoT ceiling. *Cross-schema* indirect collapses to 0.046 (Wilson 95% CI [0.016, 0.127]), worse than the no-adapter base. The strong invocation hypothesis (“adapter learns to recite its facts at any indirect question”) is empirically falsified.



Table 33: Recite-then-reason traces mixed into Stage A.

| condition                                          | direct       | indirect     |
|----------------------------------------------------|--------------|--------------|
| Stage A pilot (paraphrase + QA)                    | 1.000        | 0.158        |
| POLAR + explicit CoT prompt                        | —            | 0.296        |
| ICL upper (facts in context)                       | 0.515        | 0.304        |
| <b>Stage A + recite traces, within-schema held</b> | <b>0.985</b> | <b>0.407</b> |
| Stage A + recite traces, cross-schema held         | 0.985        | <b>0.046</b> |

Table 34: Stage C: full-parameter base meta-trained over LoRA distribution. 8 train uids, 2 held uids, 300 steps at LR 1e-5.

| split       | indirect (pre Stage C) | indirect (post Stage C)      | direct (post Stage C) |
|-------------|------------------------|------------------------------|-----------------------|
| train users | 0.147                  | <b>0.269</b>                 | 0.219 (collapsed)     |
| held users  | 0.167                  | 0.167 ( <b>no transfer</b> ) | 0.312                 |

#### A.4 Stage C meta-train (failed)

Two findings: (1) Held-user indirect lift is +0.000—the “Introspection Transfer” hypothesis fails at minimum scale. (2) Direct recall on the train users *collapses* from 1.000 to 0.22 because the base fine-tune misaligns the frozen LoRAs.

#### A.5 Diagnosis: LoRA’s global-edit failure mode

The empirical picture from User-as-LoRA is consistent: per-user LoRA contains the facts (CoT-prompted recall  $\approx$  ICL ceiling); LoRA does not spontaneously surface them on indirect queries (recall 0.150); the LoRA’s global weight perturbation actively interferes with the model’s general reasoning (5/10 users have indirect with LoRA *worse* than without). Recite-then-reason traces help within-schema but not cross-schema. Stage C meta-training over the LoRA distribution does not learn to introspect held-out adapters at the scales we tested.

These findings collectively motivate a per-user memory mechanism where (a) the gate-vs-store separation is architectural rather than learned, and (b) the per-user edit is local enough to never contaminate non-trigger forwards. Engram with surgical row insertion (main paper Section ??) provides exactly this property.

## B Appendix: Mechanistic analysis

## C Appendix: Pretraining curves

## D Appendix: Insertion strategy comparison

## E Appendix: Paraphrase generalisation

## F Appendix: Multi-domain composition

## G Appendix: Serving latency CDF

## H Appendix: Qwen-3B + RAG details

The Qwen-3B + RAG measurements in Table ?? use Qwen/Qwen2.5-3B-Instruct in bf16 with greedy decoding ( $\leq 32$  new tokens). The chat-template prompt is:

[SYSTEM]

You are answering personal-memory questions about the user. Use only




Figure 16: LogitLens KL by layer (Mini-Engram-d8 vs. base-d8). Engram converges faster at layer 3 (−3.66 KL gap), confirming Cheng et al.’s effective-deepening claim at our scale.

Figure A2: Mini-Engram pretraining curves.




Figure 17: Mini-Engram pretraining curves. Engram d8 reaches lower validation bpb than base d8 at iso-FLOPs (matching the paper’s finding at much smaller scale); engram d12 is substantially better thanks to the larger backbone and longer training.

the facts provided. Answer with the shortest possible span (typically 1-3 words). Do not explain or add any extra text.

[USER]

Facts about the user:

- <retrieved fact 1>.
- <retrieved fact 2>.
- ...

Question: <indirect question from indirect\_qa>

Answer:

Retrieval uses the same sentence-transformers/all-MiniLM-L6-v2 encoder as the RAG/MEM0/MEMMACHINE baselines of Section ??, with the per-user fact set as the index. For each indirect probe the question text is embedded and cosine-similarity top- $k$  retrieval returns the facts to put in the context block. The same 20-user/20-probe split as Table ?? is evaluated.

Two answer-match criteria are reported: *indirect\_top1* requires the first generated word to equal (case-insensitive) the first word of the gold answer; *indirect\_any* requires the gold span to occur anywhere in the generation. This matches the substring matcher used for F.

**Why oracle-top-1 underperforms RAG top-1.** A non-obvious result in Table ?? is that oracle top-1 retrieval (use `required_fact_keys` to fetch the ground-truth fact) gives 46% *indirect\_any*,

Figure A3: Insertion strategy comparison.




Figure 18: Top-1 and top-5 recall at 100 facts/user, Mini-Engram-d12. RANDOM (control) confirms our signal isn’t noise; UNEMBED\_P alone is too weak; OPT-15 helps; Joint OPT closes most of the gap to full LoRA fine-tuning at  $161\times$  less storage.




Figure 19: Paraphrase generalisation across 4 facts  $\times$  5 paraphrases each. Single-trigger insertion gets 50% top-1 for free (suffix N-gram overlap); multi-trigger insertion (insert at all 5 paraphrases) gets 100% at  $5\times$  the per-fact OPT cost.

below RAG top-1’s 55%. Inspecting the data, retrieval often returns facts that are *adjacent* to the required key in semantic space and that happen to also contain the answer surface (e.g. retrieving the “birth-year” fact when asked about “current age” covers both the year and the implicit age). RAG with  $k = 1$  thus over-samples the information needed for indirect reasoning, while strict oracle retrieval leaves the LM to derive the answer from one single-attribute fact—which Qwen-3B often refuses to do under our “shortest span” system prompt.

## I Appendix: Preliminary architectural validation

Before training Mini-Engram, we ran three feasibility probes on the randomly-initialised Engram demonstration code from ? to verify the architecture admits surgical insertion at all. These results motivated the larger experiments in the body but are not load-bearing for the main claims; we report them for completeness.

**T1 collision audit (100 synthetic users  $\times$  30 facts).** Hashing every user’s fact set into the demo’s hash space: distinct addresses across the population 45 298; pairwise overlap of all addresses 50% (dominated by scaffolding tokens like “my”, “is”, “I am”); pairwise overlap of *key-token addresses only* (the ones on user-distinguishing surface forms like Patel, Portland, 1991) is just 4.5%. For a 30-fact user the slot-occupancy fraction is 0.027%; the 1.6 M-slot demo table is 99.97% empty after one user. *Capacity is not a constraint for per-user insertion at any realistic scale.*

Figure A5: Multi-domain composition (top-1 per domain).  
Heavy degradation when domains share trigger templates.




Figure 20: Multi-domain additive composition heatmap. Each cell shows per-domain top-1 recall when  $D$  domains are stacked. Heavy degradation when domains share trigger templates (high address overlap). Disjoint-template domains (corp + user, demonstrated in main text) compose without loss.

Figure A6: Per-request latency CDF (d12, 30u×50f).




Figure 21: Per-request latency CDF on Mini-Engram-d12 (30 users  $\times$  50 facts  $\times$  600 requests). Override apply (blue) and restore (orange) are both sub-3 ms at p99; the forward pass (grey) dominates total latency (black).

**T2 surgical read/write existence proof (random-init demo).** Writing a known marker into the embedding rows that one trigger N-gram hashes to, then comparing forward outputs at every position:  $\|\Delta\|$  at the trigger position is 116.96; mean  $\|\Delta\|$  at non-trigger positions is 0.65 (a  $179.7\times$  ratio). The empirical change at the trigger position has cosine similarity 0.998 with the predicted  $W_V \cdot \text{marker}$  projection. Architecture admits surgical insertion mechanically; the trained model improves these numbers further (Figure ??).

**T3 cross-user collision rates.** Across 100 synthetic users with  $\sim 30$  facts each, cross-user hash collisions on *distinguishing* tokens (the ones that disambiguate users) occur at 4.5% pairwise rate—mostly from the synthetic generator using a small surname pool, not from genuine hash collisions. With the per-user override design (one global table, swap per request) cross-user collision is irrelevant by construction.

## J Appendix: Alternative architectures we considered

For completeness we report two alternative designs we benchmarked before settling on per-user override tables.

Table 35: Shared-table-with-salt: own-fact recall vs. leak rate. All numbers on Mini-Engram-d12 with the XL corpus.

| config                                              | own-fact top-1 | cross-user leak |
|-----------------------------------------------------|----------------|-----------------|
| no salt + UNEMBED_P (cheap, $100 \times 100$ facts) | 0.089          | 0.039           |
| salt + OPT-15 ( $10 \times 30$ facts)               | 0.11           | <b>0.000</b>    |
| salt + strong OPT-60 ( $10 \times 30$ )             | 0.46           | <b>0.000</b>    |
| salt + strong OPT-60 ( $30 \times 100$ )            | 0.345          | 0.005           |

Table 36: Insertion strategy comparison on the original 8 USER + 8 ORG benchmark, Mini-Engram-d8. RANDOM is a control; UNEMBED\_P is the closed-form pseudo-inverse; OPT is 15-step gradient on the row. Numbers are mean across the 16 facts.

| strategy              | mean $\Delta\text{logit}$ | $+\Delta\text{logit}$ | rank improved | top-1 hits  | top-5 hits  |
|-----------------------|---------------------------|-----------------------|---------------|-------------|-------------|
| RANDOM                | -4.708                    | 1/16                  | 2/16          | 0/16        | 0/16        |
| WTE                   | -0.150                    | 6/16                  | 6/16          | 0/16        | 1/16        |
| UNEMBED_P             | +1.711                    | 15/16                 | 13/16         | 1/16        | 3/16        |
| <b>OPT (15 steps)</b> | <b>+5.323</b>             | <b>16/16</b>          | <b>15/16</b>  | <b>6/16</b> | <b>7/16</b> |

**Shared table with per-user hash salt (deprecated).** In this design, every user’s facts coexist in one physical table; their hashes are XOR-salted by user-id so identical surface triggers map to disjoint rows. We benchmark on Mini-Engram-d12:

The salt design provides architectural privacy at the cost of recall: salted addresses point to untrained rows (the model never saw salt during pretraining), so the gate fires more weakly and OPT must work harder. Because per-user override tables (main body Section ??) provide stronger privacy by construction *and* use the trained address space, we recommend overrides as the production design and keep salt as a fallback when the storage backend cannot afford per-user override maps.

**100-user serving: UNEMBED\_P vs. Joint OPT.** At 100 users  $\times$  100 facts on Mini-Engram-d12@1280, the cheap (closed-form) UNEMBED\_P strategy reaches 9.2% top-1 / 23.4% top-5 at 2.3 ms apply latency. Joint OPT (1000 steps, the production setting) on the same configuration reaches **60% top-1 / 97% top-5**, matching the  $\sim 60\%/95\%$  prediction extrapolated from the 30-user data and consistent with the per-user-density curve at d12@1280 (Table ??: 67% top-1 / 99% top-5 at 100 facts/user). The serving-protocol cross-user leak under Joint OPT is 4.4% from gold-value coincidence (5 of 114 cross-user probes), versus 46% under the all-coresident protocol where every user’s overrides occupy the same table at once.

## K Appendix: Insertion strategies (full)

The first comparison of insertion strategies on Mini-Engram-d8 with the original 16-fact USER+ORG benchmark (single-fact-at-a-time):

The DUAL strategy (jointly satisfying gate-K and value-V via least-squares) was tested but *worse* than UNEMBED\_P ( $4/16 + \Delta\text{logit}$ ) — over-constraining  $e$  drives it off the trained-row distribution and the gate suppresses it. We dropped DUAL.

## L Appendix: Per-paraphrase rank table

The full per-paraphrase rank data behind Section ?? and Figure ??, single-trigger insertion on Mini-Engram-d8:

Generalisation is high when paraphrases share suffix N-grams (e.g., all “doctor” paraphrases end in “Dr.”); poor when surface form diverges (“I love the spice” vs. “My favorite spice is”). Inserting at all 5 paraphrases per fact (multi-trigger) drives every paraphrase to rank 0.

Table 37: Per-paraphrase rank of the gold token after single-trigger OPT insertion. “★” marks rank 0 (top-1).

| insertion trigger                   | query (paraphrase)                  | rank |
|-------------------------------------|-------------------------------------|------|
| My doctor’s name is Dr.             | My doctor’s name is Dr.             | 0 ★  |
|                                     | My doctor is Dr.                    | 0 ★  |
|                                     | Who is my doctor? Dr.               | 54   |
|                                     | My physician’s name is Dr.          | 0 ★  |
|                                     | My GP is Dr.                        | 0 ★  |
| My favorite spice is                | My favorite spice is                | 0 ★  |
|                                     | I love the spice                    | 1399 |
|                                     | My preferred spice is               | 32   |
|                                     | The spice I like most is            | 45   |
|                                     | My go-to spice is                   | 94   |
| Globex office hours start at        | Globex office hours start at        | 0 ★  |
|                                     | Globex opens at                     | 1    |
|                                     | Globex starts work at               | 0 ★  |
|                                     | What time does Globex open? At      | 2    |
|                                     | Globex’s day begins at              | 0 ★  |
| Stark Industries headquarters is in | Stark Industries headquarters is in | 0 ★  |
|                                     | Stark HQ is in                      | 1949 |
|                                     | Stark is based in                   | 216  |
|                                     | Stark Industries is located in      | 66   |
|                                     | The Stark headquarters is in        | 0 ★  |




Figure 22: Joint OPT convergence on Mini-Engram-d12. At 100 facts/user, loss converges to 0.8 within 1500 steps; at 1000 facts, loss saturates around 2.0 due to higher within-user fact-density interference (Section ??).

## M Appendix: Joint OPT convergence

## N Appendix: Joint OPT algorithm

1. For each fact  $f$ : compute trigger global rows  $R_f$  (deterministic from tokens) and UNEMBED\_P initial marker  $e_f^{(0)} = W_V^\dagger U_{y_f}$ .
2. Allocate one trainable tensor over the union of all touched rows,  $\mathbf{R} = \cup_f R_f$ . Initialise: each row  $\mathbf{R}[i] = \text{mean}_{f:i \in R_f} e_f^{(0)}[i]$  (average of UNEMBED\_P rows where multiple facts share an address).
3. Zero out the embedding rows in  $\mathbf{R}$ . Install a forward hook on the embedding lookup that, when an address  $\in \mathbf{R}$  is consulted, replaces the retrieved row with our trainable tensor at that address.
4. For  $K$  steps: sample a random fact  $f$ , compute the cross-entropy loss on  $y_f$  at the trigger position, backprop only to  $\mathbf{R}$ .

5. After training, write **R** into the embedding table as the user’s override map.

We use  $K = 2000$  for 100 facts and  $K = 8000$  for 1000 facts, Adam LR 0.5. Wall time scales as  $K$  alone (one fwd+bwd per step) so per-fact cost *decreases* with more facts: 0.44 s/fact at  $N=100$ , 0.23 s/fact at  $N=1000$ .

## O Appendix: Seed variance of the layered measurement

The headline F vs B numbers in Table ?? are from a single LoRA training run; for transparency, we re-ran the same configuration with a different RNG state and compared. Because the training-data ordering, LoRA-weight initialisation, and Adam-state warmup are stochastic, per-user val\_bpb  $\Delta$  varies across seeds even with identical hyperparameters.

Table 38: Three-seed variance on the layered F vs B contrast, Mini-Engram-d20,  $n=20$  users per seed, identical config across seeds (rank-64 LoRA, 1500 steps, lr 5e-4, lora\_alpha 128; per-user Engram J-OPT 1500 steps lr 0.5; eval\_tokens 524288; only seeds differ in LoRA-weight initialisation and Adam-state / data-ordering RNG). The architectural property (per-user Engram  $\Delta$ bpb at noise floor  $\sim 10^{-5}$ ) is seed-invariant; the LoRA contamination magnitude varies by  $\sim 13\%$  and F indirect\_any varies by  $\sim 10$  pp across seeds while the qualitative ranking ( $F \gg B$  on indirect,  $B \gg F$  on contamination) is preserved in every seed.

| quantity                    | seed S0        | seed S1        | seed S2        | 3-seed mean (range)   |
|-----------------------------|----------------|----------------|----------------|-----------------------|
| B $\Delta$ bpb mean         | +1.784         | +1.754         | +1.557         | +1.698 [+1.56, +1.78] |
| B $\Delta$ bpb per-user rng | [+0.44, +3.74] | [+0.73, +2.46] | [+0.69, +2.89] | —                     |
| C $\Delta$ bpb mean         | +0.000052      | +0.000049      | +0.000054      | +0.000052             |
| D $\Delta$ bpb mean         | +1.819         | +1.756         | +1.290         | +1.622 [+1.29, +1.82] |
| F $\Delta$ bpb mean         | +0.386         | +0.388         | +0.379         | +0.385 [+0.38, +0.39] |
| F indirect_any              | 44.5%          | 44.2%          | 34.8%          | 41.2% [34.8, 44.5] %  |
| B indirect_any              | 6.0%           | 8.8%           | 7.2%           | 7.3% [6.0, 8.8] %     |
| F direct top-1              | 100%           | 100%           | 100%           | 100%                  |
| B direct top-1              | 99%            | 97.5%          | 100%           | 98.8% [97.5, 100] %   |
| F worse than A              | 0/20           | 0/20           | 0/20           | <b>0/60</b>           |
| B worse than A              | 17/20          | 16/20          | 16/20          | <b>49/60 (82%)</b>    |
| D worse than A              | 18/20          | 16/20          | 16/20          | <b>50/60 (83%)</b>    |

Reading: the architectural-property comparison ( $C \Delta$ bpb  $\ll B \Delta$ bpb by  $\sim 33,000\times$ ) is seed-stable to four decimal places—every seed’s C is at the noise floor and every seed’s B is in [+1.56, +1.78]. The F vs B Pareto-domination conclusion is also seed-stable: F is never worse than the no-edit base on indirect (0/60 across seeds) while B is worse in 49/60 = 82%. The exact magnitudes vary modestly: seed-to-seed F indirect\_any spans 34.8%–44.5% (mean 41.2%), and the F vs B indirect ratio spans  $4.8\times$  (S2) to  $7.4\times$  (S0). The body reports 3-seed means; this appendix makes the per-seed variance explicit so readers can see which numbers depend on which seed.

## P Appendix: Reproducing the paper

```
git clone https://github.com/bojieli/user-as-engram
cd user-as-engram/nanochat
uv venv && source .venv/bin/activate && uv sync --extra gpu
export NANOCHAT_BASE_DIR=$(pwd)/../nanochat_base
export CKPT=$NANOCHAT_BASE_DIR/engram_runs/engram_d12

# Pretrain Mini-Engram-d12 (~3h on a single GPU)
python -m scripts.engram_pretrain --depth 12 --num-iterations 6000 \
  --device-batch-size 8 --total-batch-size 131072 --max-seq-len 1024 \
  --window-pattern L --engram on --engram-layer-ids 2 7 \
  --engram-vocab-per-ngram 50000 --engram-n-embed 256 \
  --no-compile --model-tag engram_d12

# Build corpora
python -m scripts.build_corpus_xxl

# Run all evaluations
```

```
python -m scripts.scalability_benchmark --ckpt-dir $CKPT
python -m scripts.joint_opt             --ckpt-dir $CKPT --n-facts 100
python -m scripts.multifact_lora         --ckpt-dir $CKPT --n-facts 100
python -m scripts.eval_serving           --ckpt-dir $CKPT --n-users 30
```